

MAS DSE 201 Homework: 201Cats

Milestone IV [due March 18 midnight]

Your next goal will be to improve the performance of the “My Kind of cats – with preference” option by appropriate precomputations. It is understood that the maintenance of the precomputed tables will lead to some slowdown while viewers interact. Your precomputation choices must be such that the precomputations that you introduce collectively save more time to this option than they cost to maintain. Precomputed tables will also benefit from creating certain indices on them. Build any beneficial indices. Calibrate your solution against cold runs of the activity script provided.

Submit the following.

- The precomputed tables you created: `CREATE TABLE` statement and query that initially loaded it.
- Your new “My Kind of cats – with preference” query, which makes the best use of the precomputed tables you choose.
- The indices you created (`CREATE INDEX` scripts)

201 Cats - Summary of Database Size

| Table Name | Number of Tuples |
|--------------------------------------|------------------|
| users | 10,000 |
| sessions (not used in queries) | 5 |
| videos | 25,000 |
| friends | 150,000 |
| suggestions (not used in queries) | 5 |
| likes | 100,000 |
| watched (not used in queries) | 5 |

201 Cats - Pre-Computed Tables Created

| Approach | Table Name | Run Time (without Indices) | Run Time (with Indices) |
|----------|---------------------|----------------------------|-------------------------|
| #1 | calculate_userX | 124 msec | 107 msec |
| | calculate_userY | 10 min 40 secs | 12 min 17 secs |
| #2 | calculate_logcosine | 4 min 38 secs | 4 min 13 secs |

Screenshots for Approach #1

```
-- calculate and select the vector for parameterized user X
CREATE TABLE calculate_userX AS
select v.video_id, u.user_id, (case when l.like_id is null then 0 else 1 end) as liked
from videos v
cross join users u
left join likes l
on u.user_id = l.user_id and v.video_id = l.video_id
where u.user_id = 3 -- specify user X
;
```

```
-- calculate and select the vector for all users Y
CREATE TABLE calculate_userY AS
select v.video_id, u.user_id, (case when l.like_id is null then 0 else 1 end) as liked
from videos v
cross join users u
left join likes l
on u.user_id = l.user_id and v.video_id = l.video_id
where u.user_id != 3 -- exclude user X
;
```

Run Time for *calculate_userX* - without Indices (left) and with indices (right)

| | | |
|---|----------------------|----------|
| 3/16/2020 9:23:49 PM | 25,000 | 124 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| <pre>-- calculate and select the vector for parameterized user X CREATE TABLE calculate_userX AS select v.video_id, u.user_id, (case when l.like_id is null</pre> | | |

| | | |
|---|----------------------|----------|
| 3/16/2020 8:54:37 PM | 25,000 | 107 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| <pre>-- calculate and select the vector for parameterized user X CREATE TABLE calculate_userX AS select v.video_id, u.user_id, (case when l.like_id is null t</pre> | | |

Run Time for *calculate_userY* - without Indices (left) and with indices (right)

| | | |
|---|----------------------|----------------|
| 3/16/2020 9:26:12 PM | 249,975,000 | 10 min 40 secs |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| <pre>-- calculate and select the vector for all users Y CREATE TABLE calculate_userY AS select v.video_id, u.user_id, (case when l.like_id is null then</pre> | | |

| | | |
|---|----------------------|----------------|
| 3/16/2020 8:54:48 PM | 249,975,000 | 12 min 17 secs |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| <pre>-- calculate and select the vector for all users Y CREATE TABLE calculate_userY AS select v.video_id, u.user_id, (case when l.like_id is null then</pre> | | |

Screenshots for Approach #2

```

-- calculate the log cosine for user X and all users Y
CREATE TABLE calculate_logcosine AS
-- calculate inner product and log cosine
select user_y, log(sum(product)+1) as lc
from(
  -- calculate and select multiplication from inner product
  select t1.video_id, t1.user_id as user_x, t2.user_id as user_y, t1.liked as liked_x, t2.liked
    as liked_y, (t1.liked * t2.liked) as product
  from (
    -- calculate and select the vector for parameterized user X
    select v.video_id, u.user_id, (case when l.like_id is null then 0 else 1 end) as liked
    from videos v
    cross join users u
    left join likes l
    on u.user_id = l.user_id and v.video_id = l.video_id
    where u.user_id = 3 -- specify user X
  ) as t1
  left join (
    -- calculate and select the vector for all users Y
    select v.video_id, u.user_id, (case when l.like_id is null then 0 else 1 end) as liked
    from videos v
    cross join users u
    left join likes l
    on u.user_id = l.user_id and v.video_id = l.video_id
    where u.user_id != 3 -- exclude user X
  ) as t2
  on t1.video_id = t2.video_id
) as t3
group by user_y
;

```

Run Time for *calculate_logcosine* – without Indices (left) and with indices (right)

| | | |
|--|----------------------|---------------|
| 3/16/2020 9:55:34 PM | 9,999 | 4 min 38 secs |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- calculate the log cosine for user X and all users Y | | |

| | | |
|--|----------------------|---------------|
| 3/16/2020 10:09:15 PM | 9,999 | 4 min 13 secs |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- calculate the log cosine for user X and all users Y | | |
| CREATE TABLE calculate_logcosine AS | | |

201 Cats - New “My Kind of cats – with preference” queries

| Query Type | Run Time (with indices) | Run Time (without indices) |
|---|----------------------------|-------------------------------|
| Original | 4 min 19 secs | 4 min 19 secs |
| With Pre-Computed Tables (Approach #1) | 2 min 24 secs | 2 min 23 secs |
| With Pre-Computed Tables (Approach #2) | 361 msec | 351 msec |

Original

```
-- My kind of cats - with preference
select video_id, "Video Name", "Sum Weighted Likes"
from (
  -- select videos with ranked log cosine
  select l.video_id, v.name as "Video Name", cast(sum(t4.lc) as decimal(36,4)) as "Sum Weighted Likes",
  dense_rank() over(order by sum(t4.lc) desc) as rn
  from likes l
  join(
    -- calculate inner product and log cosine
    select user_y, log(sum(product)+1) as lc
    from(
      -- calculate and select multiplication from inner product
      select t1.video_id, t1.user_id as user_x, t2.user_id as user_y, t1.liked as liked_x, t2.liked
      as liked_y, (t1.liked * t2.liked) as product
      from (
        -- calculate and select the vector for parameterized user X
        select v.video_id, u.user_id, (case when l.like_id is null then 0 else 1 end) as liked
        from videos v
        cross join users u
        left join likes l
        on u.user_id = l.user_id and v.video_id = l.video_id
        where u.user_id = 3 -- specify user X
      ) as t1
      left join (
        -- calculate and select the vector for all users Y
        select v.video_id, u.user_id, (case when l.like_id is null then 0 else 1 end) as liked
        from videos v
        cross join users u
        left join likes l
        on u.user_id = l.user_id and v.video_id = l.video_id
        where u.user_id != 3 -- exclude user X
      ) as t2
      on t1.video_id = t2.video_id
    ) as t3
    group by user_y
  ) as t4
  on l.user_id = t4.user_y
  join videos v
  on l.video_id = v.video_id
  group by l.video_id, v.name
) as t5
-- filter for top videos
where rn <= 10
;
```

| Date | Rows Affected | Duration |
|-----------------------|---------------|---------------|
| 3/16/2020 10:35:31 PM | 24,548 | 4 min 19 secs |

Copy Copy to Query Editor

-- My kind of cats - with preference

Approach #1

```
-- Cats: New Query with Precomputed Tables
-- My kind of cats - with preference
select video_id, "Video Name", "Sum Weighted Likes"
from (
  -- select videos with ranked log cosine
  select l.video_id, v.name as "Video Name", cast(sum(t4.lc) as decimal(36,4)) as "Sum Weighted Likes",
  dense_rank() over(order by sum(t4.lc) desc) as rn
  from likes l
  join(
    -- calculate inner product and log cosine
    select user_y, log(sum(product)+1) as lc
    from(
      -- calculate and select multiplication from inner product
      select t1.video_id, t1.user_id as user_x, t2.user_id as user_y, t1.liked as liked_x, t2.liked
      as liked_y, (t1.liked * t2.liked) as product
      from calculate_userX t1 -- use precomputed table
      left join calculate_userY t2 -- use precomputed table
      on t1.video_id = t2.video_id
    ) as t3
    group by user_y
  ) as t4
  on l.user_id = t4.user_y
  join videos v
  on l.video_id = v.video_id
  group by l.video_id, v.name
) as t5
-- filter for top videos
where rn <= 10
;
```

| Date | Rows Affected | Duration |
|-----------------------|---------------|---------------|
| 3/16/2020 10:40:57 PM | 24,548 | 2 min 24 secs |

Copy Copy to Query Editor

-- Cats: New Query with Precomputed Tables
-- My kind of cats - with preference

Approach #2

```
-- Cats: New Query with Precomputed Tables
-- My kind of cats - with preference
select video_id, "Video Name", "Sum Weighted Likes"
from (
  -- select videos with ranked log cosine
  select l.video_id, v.name as "Video Name", cast(sum(t4.lc) as decimal(36,4)) as "Sum Weighted Likes",
  dense_rank() over(order by sum(t4.lc) desc) as rn
  from likes l
  join calculate_logcosine t4
  on l.user_id = t4.user_y
  join videos v
  on l.video_id = v.video_id
  group by l.video_id, v.name
) as t5
-- filter for top videos
where rn <= 10
;
```

| Date | Rows Affected | Duration |
|-----------------------|---------------|----------|
| 3/16/2020 10:48:00 PM | 24,548 | 361 msec |

Copy Copy to Query Editor

-- Cats: New Query with Precomputed Tables
-- My kind of cats - with preference

201 Cats - Indices Created

| Table | Column | Used In Clause | SQL Statement |
|--------|---------------------------|-------------------------------|--|
| videos | name | GROUP BY | CREATE INDEX videos_name_idx ON videos(name); |
| likes | video_id (Foreign Key) | LEFT JOIN JOIN GROUP BY | CREATE INDEX likes_videoid_idx ON likes(video_id); |
| likes | user_id (Foreign Key) | LEFT JOIN JOIN | CREATE INDEX likes_userid_idx ON likes(user_id); |

201 Cats – Summary of Observations

- Two approaches were taken to implement pre-computed tables:
 - For approach #1: a table is created for each vector X and Y, where X is the user specified and Y are all the other users. This approach is useful because it takes care of all the expensive computations happening during the cross join between tables "users" and "likes", however, generating the pre-computed table "calculate_userY" requires a significant amount of memory and run time with 249,975,000 tuples and more than 10 minutes to execute (in this database tested).
 - For approach #2: a table is created to take care of similar computations mentioned on approach #1 but to also to handle joining the tables for vectors X and Y, plus performing the vector inner product and the log cosine calculation.
- Indices
 - On pre-computed tables: Indices seem to be benefiting tables "calculate_userX" and "calculate_logcosine", but affecting table "calculate_userY" from 10min 40sec to 12min 17sec in run time. As mentioned before, table "calculate_userY" is quite large and is possible that memory access and writing makes the run time slower.
 - On queries: Using indices to run the queries, for both pre-computed tables and without, does not reflect an impact in run time; on milestone 3 I previously showed the same result for the query without pre-computed tables. However, what I believe is interesting on this homework is that indices are probably not helping the query as a whole but when you create a pre-computed table on a smaller segment this one can reflect an improvement in run time.
- The Best Approach/Conclusion
 - Approach #2 is definitely performing better than #1. Creating the table "calculate_logcosine" will take less than half of the time needed to create "calculate_userY" alone (for the reasons mentioned previously). Another step that is very expensive on approach #1 is the left join between "calculate_userX" and "calculate_userY", considering that the latter table will always be quite large in reference to the former.
 - Approach #2 is also a better option than the original query, when comparing the best scenarios running it collectively (pre-computed table + query) takes less than 4 min 14 sec while the original query takes 4 min 19 sec.

MAS DSE 201 Homework: Sales Cube

Milestone IV [due March 18 midnight]

Your next goal will be to improve the performance of Query 6 by appropriate precomputations. It is understood that the maintenance of the precomputed table(s) will lead to some slowdown while viewers interact. Your precomputation choices must be such that the precomputations that you introduce collectively save more time to this option than they cost to maintain. Precomputed tables will also benefit from creating certain indices on them. Build any beneficial indices. Calibrate your solution against cold runs of the activity script.

Submit the following.

- The precomputed tables you created: `CREATE TABLE` statement and query that initially loaded it.
- Your new Query 6, which makes the best use of the precomputed tables you chose.
- The indices you created (`CREATE INDEX` scripts)

Sales Cube - Summary of Database Size

| Table Name | Number of Tuples |
|------------|------------------|
| states | 5,000 |
| customers | 700,000 |
| categories | 900,000 |
| products | 800,000 |
| sales | 1,000,000 |

Sales Cube - Pre-Computed Tables Created

| Table Name | Run Time (without Indices) | Run Time (with Indices) |
|-----------------------------|----------------------------|-------------------------|
| top20_customers | 5 secs 765 msec | 6 secs 420 msec |
| top20_categories | 4 secs 666 msec | 4 secs 676 msec |
| sales_customerandcategories | 8 secs 803 msec | 6 secs 875 msec |

```
-- select top 20 for customers
CREATE TABLE top20_customers AS
  select *
  from (
    select c.customer_id, c.first_name, c.last_name, dense_rank() over(order by sum(s.price_paid) desc) as rn
    from sales s
    join customers c
    on s.customer_id = c.customer_id
    group by c.customer_id
  ) as temp_t
  where rn <= 20
;
```

Run Time for *top20_customers* – without Indices (left) and with indices (right)

| | | |
|--------------------------------|----------------------|-----------------|
| 3/17/2020 12:33:34 AM | 22 | 5 secs 765 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- select top 20 for customers | | |

| | | |
|--------------------------------|----------------------|-----------------|
| 3/17/2020 12:43:46 AM | 22 | 6 secs 420 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- select top 20 for customers | | |

```
-- select top 20 for categories
CREATE TABLE top20_categories AS
select *
from (
    select cat.category_id, cat.name, dense_rank() over(order by sum(s.price_paid) desc) as rn
    from sales s
    join products p
    on s.product_id = p.product_id
    join categories cat
    on cat.category_id = p.category_id
    group by cat.category_id, cat.name
) as temp_t
where rn <= 20
;
```

Run Time for *top20_categories* – without Indices (left) and with indices (right)

| | | |
|---------------------------------|----------------------|-----------------|
| 3/17/2020 12:33:20 AM | 20 | 4 secs 666 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- select top 20 for categories | | |

| | | |
|---------------------------------|----------------------|-----------------|
| 3/17/2020 12:43:23 AM | 20 | 4 secs 676 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- select top 20 for categories | | |

```
-- joining tables for sales, products, and categories
CREATE TABLE sales_customerandcategories AS
select s.price_paid, s.customer_id, cat.category_id, s.quantity
from sales s
join products p
on s.product_id = p.product_id
join categories cat
on cat.category_id = p.category_id
;
```

Run Time for *sales_customerandcategories* – without Indices (left) and with indices (right)

| | | |
|---|----------------------|-----------------|
| 3/17/2020 12:33:05 AM | 1,500,000 | 8 secs 803 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- joining tables for sales, products, and categories | | |
| CREATE TABLE sales_customerandcategories AS | | |

| | | |
|---|----------------------|-----------------|
| 3/17/2020 12:42:58 AM | 1,500,000 | 6 secs 875 msec |
| Date | Rows Affected | Duration |
| Copy | Copy to Query Editor | |
| -- joining tables for sales, products, and categories | | |

Sales Cube – Queries for #6

```
-- 6. Tuples for each of the top 20 product categories and top 20 customers
select t5.category_id, t5.name as "top category", t5.customer_id, t5.first_name
as "top customer first name", t5.last_name as "top customer last name",
coalesce(sum(s.quantity),0) as "quantity sold", coalesce(sum(s.price_paid),0) as "dollar value"
from sales s
join products p
on s.product_id = p.product_id
join categories cat
on cat.category_id = p.category_id
right join (
  select * from (
    -- select top 20 for categories
    select *
    from (
      select cat.category_id, cat.name, dense_rank() over(order by sum(s.price_paid) desc) as rn
      from sales s
      join products p
      on s.product_id = p.product_id
      join categories cat
      on cat.category_id = p.category_id
      group by cat.category_id, cat.name
    ) as t1
    where rn <= 20
    ) as t2
    cross join ( -- find all possible combinations ~400 with cross join
    -- select top 20 for customers
    select *
    from (
      select c.customer_id, c.first_name, c.last_name, dense_rank() over(order by sum(s.price_paid) desc) as rn
      from sales s
      join customers c
      on s.customer_id = c.customer_id
      group by c.customer_id
    ) as t3
    where rn <= 20
    ) as t4
  ) as t5
-- join on categories and customers only
on cat.category_id = t5.category_id and s.customer_id = t5.customer_id
group by t5.category_id, t5.name, t5.customer_id, t5.first_name, t5.last_name
order by "dollar value" desc
;
```

Original

| Query Type | Run Time (with indices) | Run Time (without indices) |
|--------------------------|----------------------------|-------------------------------|
| Original | 16 secs 545 msec | 14 secs 986 msec |
| With Pre-Computed Tables | 813 msec | 802 msec |

3/17/2020 12:58:06 AM

440

16 secs 545 msec

Date

Rows Affected

Duration

Copy

Copy to Query Editor

```
-- 6. Tuples for each of the top 20 product categories and top 20 cus
select t5.category_id, t5.name as "top category", t5.customer_id, t5
```

3/17/2020 12:59:45 AM

440

813 msec

Date

Rows Affected

Duration

Copy

Copy to Query Editor

```
-- Sales: New Query with Precomputed Tables
```

```
-- 6. Tuples for each of the top 20 product categories and top
```

```
-- Sales: New Query with Precomputed Tables
-- 6. Tuples for each of the top 20 product categories and top 20 customers
select t5.category_id, t5.name as "top category", t5.customer_id, t5.first_name
as "top customer first name", t5.last_name as "top customer last name",
coalesce(sum(s_cc.quantity),0) as "quantity sold", coalesce(sum(s_cc.price_paid),0) as "dollar value"
from sales_customerandcategories s_cc -- use precomputed table
right join (
  select *
  from top20_categories -- use precomputed table
  cross join top20_customers -- use precomputed table
) as t5
-- join on categories and customers only
on s_cc.category_id = t5.category_id and s_cc.customer_id = t5.customer_id
group by t5.category_id, t5.name, t5.customer_id, t5.first_name, t5.last_name
order by "dollar value" desc
;
```

With Pre-Computed
Tables

Sales Cube - Indices Created

| Table | Column | Used In Clause | SQL Statement |
|------------|------------------------------|----------------|--|
| sales | customer_id (Foreign Key) | JOIN | CREATE INDEX sales_customerid_idx ON sales(customer_id); |
| sales | product_id (Foreign Key) | JOIN | CREATE INDEX sales_productid_idx ON sales(product_id); |
| products | category_id (Foreign Key) | JOIN | CREATE INDEX products_categoryid_idx ON products(category_id); |
| categories | name | GROUP BY | CREATE INDEX categories_name_idx ON categories(name); |

Sales Cube – Summary of Observations

- Pre-Computed Tables
 - “top20_customers” table: Finds the top 20 ranked customers based on the price paid. It computes the join between sales and customers, in addition to the rank. This is a nice summary for anyone that might be interested in who are the customers spending the most money in the organization's products or services.
 - “top20_categories” table: Finds the top 20 ranked categories based on the price paid. It computes the three joins among tables sales, products, and categories; in addition to the rank. This is also a nice summary for anyone that might be interested in maybe understanding what categories are bringing the most revenue or perhaps are the most popular among customers.
 - “sales_customerandcategories” table: Aggregates all the "sales" with the "product" and "categories" table information. This pre-computed table seems very useful for users that query the "sales" table frequently but also want information coming from the "products" or "categories" tables, therefore, instead of joining all three tables in every single query (for each user) we can have an aggregation that becomes computationally cheaper in the long run.
- Indices
 - On Pre-Computed Tables: Indices seem to only be beneficial for table "sales_customerandcategories" by lowering the run time from 8secs 803msec to 6secs 875msec, while "top20_customers" decreased in performance by almost 1 second and "top20_categories" observed no change. Is likely that "sales_customerandcategories" table is benefiting the most from indices because the number of tuples (1,500,000) visited in memory is many orders of magnitude higher as compared to the other top 20 tables (with ~20 tuples).
 - On queries: There is no run time improvement when testing the queries with indices. The original query (no pre-computed tables) increased from 14secs 986msecs to 16secs 545msec when indices were applied, and the query with pre-computed tables also increased from 802msec to 813msec; the latter one is probably within variation (~10msec change).

- Conclusion
 - The cost to generate the query with pre-computed tables (including indices) is about 20 seconds, this is almost 4 seconds slower than the original query with a run time of 16secs 545msec. Given that pre-computed tables are giving worst performance, this is one of those cases where possible applications have to be considered for the tables generated, and how they bring utility to the end users. As mentioned above, someone may be particularly interested in just getting the top 20 customers or categories, or maybe query a single table to get all the sales information. All and all, if the needs of the data ecosystem consider more flexibility for data availability and cheaper performance over the long run it would make sense to implement pre-computed tables, but if optimizing the whole query is the main goal a better option would be to simply implement indices or perhaps to think about other possible optimization alternatives (i.e. re-writing the query or re-designing the schema).