

# Final - Raul G. Martinez (PID: A12461871)

## DSE 220: Machine Learning

Due Date: 06/12 11:59 PM

In [ ]:

### Report Section

#### Abstract

This report aims to provide a Machine Learning approach to predict how useful an Amazon Review is going to be, the data used in this analysis contains a total of 200,000 reviews in the years of 2003 to 2014 from a total of 39,249 reviewers and 19,913 products. Natural language processing techniques were used to learn new features from the review body text and summary text data, in addition to feature engineering in many other numerical features. Then, a Random Forest Regressor Model was implemented to predict the ratio between the number of helpful votes with the total number of votes for each review; train and validation split was done 80% and 20% respectively. After generating a baseline model with three fundamental features (itemID, reviewerID, and outOf\_feature), model optimization was achieved with feature selection in two ways. First, each new feature was tested individually along with the baseline features and their model performance was evaluated with MAE (Mean Absolute Error) and MSE (Mean Squared Error), and any feature performing worst than the baseline was excluded. Second, an ablation study was performed by removing one feature at time, starting with the highest MAE, and the group of features yielding the lowest MAE were selected for downstream optimization. Next, hyper-parameter tuning was performed on max\_depth and n\_estimators parameters and the model with the lowest MAE was again selected. After final model selection, the confusion matrix was computed to evaluate the model's ability to classify the number of helpful votes; precision and recal metrics were observed in detail. All in all, this analysis and implementation emphasizes on feature engineering and feature selection rather than model or hyper-parameter optimization. A total of 36 new features were generated, tested, and selected to provide an automated optimal performance where the best MAE was found to be 0.15479 for the train dataset and 0.16357 for 60% of the test dataset.

#### Data Preparation and Pruning

The training data consists of 200,000 entries inside the compressed file 'train.json.gz', while the test data is stored also in a compressed file named 'test\_Helpful.json.gz' and it contains 14,000 entries. Both datasets contain a total of 12 raw features summarized in the table below (observations are based on training data):

Raw Feature	Observations
categoryID	unique values are 0, 1, 2, 3, and 4

Raw Feature		Observations
categories	there are 1042 unique categories (i.e. 'Active Hoodies') forming 1847 unique lists, each review has a collection of lists	
itemID		there are a total of 19,913 unique items
reviewerID		there are a total of 39,249 unique reviewers
rating	values assigned are 1, 2, 3, 4, and 5. Their frequency increases monotonically with the number, 1 being the less frequent	
reviewText		the min number of characters is 0 and the max is 22,646
reviewHash		each review has a unique hash ID
reviewTime		time goes from 2003 to 2014, the most reviews are observed in 2013 and 2014
summary		the min number of characters is 1 and the max is 201
unixReviewTime		there are 2,532 unique times
helpful	it contains a dictionary with 'outOf' and 'nHelpful'. Votes 0 and 1 comprise of 68.5% and 14.3% from the total number of reviews	
price		62.9% of the reviews have missing values, they were filled with -999 in this analysis

Raw features were read in a dataframe for the train (raw) dataset (see table 1 in the report section) then split into train and validation, 80% and 20% respectively, in order to validate feature selection and model optimization. The newly defined train data was then pruned to only include reviews where the number of votes (helpful - outOf) is greater than zero, therefore, the original train data (after the split) was used to compute performance metrics while the newly filtered data was only used to train the model. This is because the entries equal to zero are always predicted zero by the regression model from multiplying predicted label times the known 'outOf' feature (with value zero). Better performance was found empirically when implementing this approach, it basically allows the model to learn on the remaining 31.5% of the data where the predictions are not given or easily predicted.

Lastly, the labels are extracted from the 'helpful' column and are defined as the ratio between 'nHelpful' and 'outOf'. Divisio by zero is prevented with the approach mentioned above of removing entries with votes equal to zero.

## Feature Engineering

A total of 36 new features were generated as candidates to be included in the final model for test data predictions. The table below summarizes the new feature names (when applicable), the raw features used to generate it (when applicable), and lastly some description about it. Positive and negative word lists are taken from reference 1 and 2 below, and the list of stop words was downloaded from the NLTK library on python. A couple features were also inspired from reference 3.

Feature Name	Features Used From Raw Data	Description
itemID	NA	<b>Baseline Feature</b> - ID for each product or item
reviewerID	NA	<b>Baseline Feature</b> - ID for each reviewer or user
outOf_feature	helpful	<b>Baseline Feature</b> - Number of votes
categoryID	NA	IDs for the collection of category lists, the unique IDs are 0, 1, 2, 3, and 4

Feature Name	Features Used From Raw Data		Description
categories_count	categories		Number collection of category lists inside 'categories' column
category_numtrans	categories	Each category list gets a unique ID, then numeric transformation is applied by summation of IDs	
rating	NA		Rating for the given review
rating_deviation	rating		Deviation from the mean rating for each product
itemID_helpfulRate	itemID, helpful		Helpful rate calculated for each item on labeled data
itemID_numReviews	itemID		Number of reviews for each product
reviewerID_helpfulRate	reviewerID, helpful		Helpful rate calculated for each reviewer on labeled data
reviewerID_numReviews	reviewerID		Number of reviews for each reviewer
price	NA	Price for the item being reviewed, missing values are assigned the value -999	
reviewText_count_words (or summary)	reviewText, summary		Count for the number of words, stop words removed
reviewText_posWordCount (or summary)	reviewText, summary	Count for the number of positive words found, with pre-defined list of positive words	
reviewText_negWordCount (or summary)	reviewText, summary	Count for the number of negative words found, with pre-defined list of negative words	
reviewText_posWordRate (or summary)	reviewText, summary	Positive word rate for the text, with pre-defined list of positive words	
reviewText_negWordRate (or summary)	reviewText, summary	Negative word rate for the text, with pre-defined list of negative words	
reviewText_count_char (or summary)	reviewText, summary		Count for the number of characters
reviewText_count_punctu (or summary)	reviewText, summary		Count for the number of punctuation symbols
reviewText_count_firstCapital (or summary)	reviewText, summary		Count for the number of words where the first letter is capital
reviewText_avgWordLength (or summary)	reviewText, summary		Average word length for the text, stop words removed
reviewText_capitalwords (or summary)	reviewText, summary		Count for the number of capital words in the text
reviewText_ExclQue_countchar (or summary)	reviewText, summary		Count for the number of exclamation and question symbols in the text
reviewText_PunctChar_ratio (or summary)	reviewText, summary		Ratio for the count of punctuation symbols to characters in the text
summary_reviewText_charRatio	reviewText, summary	Ratio for the count of characters present in summary text by reviewText	
summary_reviewText_wordsRatio	reviewText, summary	Ratio for the count of words present in summary text by reviewText	
unixReviewTime	NA		Time elapsed in Unix time before the review was posted
unixReviewTime_delta_firstreview	unixReviewTime		Time elapsed in Unix time from first review for each item
unixReviewTime_delta_lastreview	unixReviewTime		Time elapsed in Unix time from the last review for each item
votes_time	helpful, unixReviewTime		Ratio between the number of votes and the time elapsed in Unix time

# Model Selection

Multiple regression models were evaluated including Linear Regression, Multi-Later Perceptron Regressor, Decision Trees, and Random Forest Regressor from different sklearn libraries. After parameter tunning and different feature testing, the top performer and simpler to implement turned out to be Random Forest Regressor; data not shown in this report. The model is trained to predict a regressor label, which is the ratio between 'nHelpful' by 'outOf', for every individual review entry. The regressor label was then transformed to a classification label by multiplying it with the known 'outOf' values and rounding to the nearest integer; rounding was shown to improve MAE consistently (data not shown in the report). Accuracy was then calculated for the classification and only reported in the data table 2 from the report section. In summary, a simple approach was selected for model selection in this analysis in order to spend more time with feature engineering and feature selection, after all, experimenting with the features showed higher contribution for error predictions MAE and MSE as opposed to model optimization.

# Feature Selection

A two step approach was taken to subselect features, both used random forest models with hyper-parameters 'max\_depth'=10 and 'n\_estimators'=100; these were found to be relatively stable and fast to test empirically. The first approach compares the performance of a baseline model with only three features (itemID, reviewerID, and outOf\_feature) with other 4 feature models by combining the baseline features and adding one of the 40 feature candidates (new and existing features) at a time. The top performing models are then ranked according to lower MAE (ascending) on the validation dataset and the features whose models performed lower than the baseline model are eliminated. Table 2 in the code section illustrates the approach, in this example a total of 7 features were eliminated and 33 selected for downstream selection. Next, the second approach aims to continue feature subselection by performing an ablation study, where features are eliminated one-by-one and model performance is measured. The features previously selected on baseline comparison are ranked by MAE (descending) and one feature is eliminated at a time, starting with no features eliminated to the first with the highest MAE and further continuing to eliminate all features. Again, the criteria to select a group of features is the test case with the lowest MAE on the validation dataset. Figure 1 shows scatter plots for MAE and MSE, both error metrics trend similarly and also clearly illustrate how model performance decreases as the number of eliminated features increases, the lowest MAE for this datset occured when 11 were removed, therefore a total of 22 features were finally selected for downstream model evaluation.

# Model Optimization and Evaluation

Many features were generated and then narrowed down to only subselect the ones with optimal performance. After selecting the final features, the same model used for testing above was then optimized by tuning the hyper-parameters 'max\_depth' and 'n\_estimators'; these parameters were found to be the most significant across 6 different ones tested by performing an expensive Randomized Search Cross-Validation with sklearn model\_selection library (not shown in this report). The values 10 and 12 are tested for max\_depth, and the values 100, 200, 400, and 600 for n\_estimators. By testing every possible scenario, the optimal combination was found by taking the one with the lowest MAE on the validation dataset. To illustrate this approach, a printout from the model is shown in the 'Hyperparameter Tuning' part from the code section, where the optimal parameters were found to be 12 for 'max\_depth' and 200 for 'n\_estimators' with a score of 0.172575 for MAE on the validation dataset.

To further evaluate the model's ability to classify, the confusion matrix was computed for the optimal predictions on the validation dataset. More evaluation metrics for the classes were also calculated such as True Positive Rate (recall), True Negative Rate, Precision, False Positive Rate, False Negative Rate, and Accuracy. Table 3 in the code section summarizes the results for these metrics for a total of 122 different classes. Is important to note that some labels show NaN values as the prediction rate, since not all predicted labels exist in the dataset and are marked as False Positives, also how accuracy is not a good metric for this analysis as a value of greater than 0.9 for every feature does not seem to be real. Conversely, when looking at precision and recall (see Figure 2) there is something interesting happening, it basically shows how the performance of both metrics is well correlated with label value magnitudes; or possibly label ranges. For instance, labels with values less than 2 have really good precision and recall (with greater than 0.7) and it then decreases to approximately 0.1 to 0.5 for labels with values

around 3 to 50, and finally almost every label greater than 50 has precision and recall of zero. This behavior is not very surprising and somewhat expected due to the nature of the predictions made in the regression model and data inconsistency for the distribution of number of votes in the training set; for example, larger number of votes for each review are significantly less common, and to put it in perspective there are only around 350 out of 200,000 reviews where the number of helpful votes is greater than 50. Additionally, the model is very good at predicting smaller values because of the rounding approach used for integers which basically reduces the error on the prediction, and again, performance progressively reduces with medium size values and it finally completely breaks down for larger value labels and is no longer useful. In summary, this model has the potential to generate good predictions when the number of votes is very small, however for predicting larger ranges such as 0 to 384 votes it will likely face limitations.

## Conclusion

To summarize, a random forest regressor model was presented in this report to predict the number of helpful votes for Amazon reviews. A total of 36 features were learned, tested, and subselected by comparing model performance with baseline features and by performing an ablation study on a selected group of features; as a matter of fact, most of the time spent in this analysis was around the former two steps along with natural language processing, feature engineering, and data exploration. Next, model optimization was done with hyper-parameter tuning for two variables which are 'max\_depth' and 'n\_estimators'. And lastly, performance metrics such as precision and recall were used to assess model performance on the validation dataset. The final model obtained an MAE of 0.15479 for the training dataset and for 60% of the testing dataset the MAE was found to be 0.16357, the error for the remaining 40% of the test dataset is not reported in this analysis.

## Future Steps

The approach presented above revealed a significant amount of new features and information from the raw data, however, the results seem to show that a single Random Forest Regressor is not sufficient for capturing most of the patterns in the data. Therefore, as next steps it would be useful to use multiple models, either implemented individually or together as an averaged prediction, that could possibly capture information better whether is through classification or regression. For example, a logistic regression model could be used to classify binary labels that occur very frequently such as 1 and 2 (not zero because it stays the same). Similarly, other models could be used to predict different ranges for the number of total votes such as 5 to 10 or even for greater than 50, depending on how they perform during testing. Lastly, the use of ensemble methods could also be considered for model optimization, some techniques widely used are bagging and boosting.

## References

1. Mingqing Hu and Bing Liu. "Mining and Summarizing Customer Reviews." Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD-2004), Aug 22-25, 2004, Seattle, Washington, USA.
2. Bing Liu, Mingqing Hu and Junsheng Cheng. "Opinion Observer: Analyzing and Comparing Opinions on the Web." Proceedings of the 14th International World Wide Web conference (WWW-2005), May 10-14, 2005, Chiba, Japan.
3. Song, Xia. "Predict Amazon Review Helpfulness With Xgboost, Neural Network, and LSTM Neural Network." Medium, Medium, 11 Aug. 2019, [medium.com/@songxia.sophia/predict-amazon-review-helpfulness-wihtxgboost-neural-network-and-lstm-neural-network-837a1da44f49](https://medium.com/@songxia.sophia/predict-amazon-review-helpfulness-wihtxgboost-neural-network-and-lstm-neural-network-837a1da44f49).

# Code Section

```
In [1]: %%javascript
IPython.OutputArea.prototype._should_scroll = function(lines) {
    return false;
}
```

```
In [2]: # import libraries

import pandas as pd
import numpy as np
pd.options.mode.chained_assignment = None # default='warn'
import gzip
from collections import defaultdict
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, accuracy_score, confusion_matrix
from sklearn.ensemble import RandomForestRegressor
import string
import nltk
from nltk.corpus import stopwords
import os
from sklearn.model_selection import train_test_split
from itertools import compress
import matplotlib.pyplot as plt
import itertools
```

## Define Functions

```
In [3]: # parse raw data

"""
These functions parse data from compressed files.
"""

def readGz(f):
    for l in gzip.open(f):
        yield eval(l)

def parse(path):
    g = gzip.open(path, 'rb')
    for l in g:
        yield eval(l)

def getDF(path):
    i = 0
    df = {}
    for d in parse(path):
        df[i] = d
        i += 1
    return pd.DataFrame.from_dict(df, orient='index')
```

```
In [4]: def GetPosNegWords_count_Rate(text):

    """
    This function finds the positive and negative word rates for the text by:
        - using pre-defined lists of positive and negative words
        - making all words lower-case
        - removes blank spaces and punctuations
    """

    # import positive and negative word lists, define as set for higher efficiency
    posWords_list = set([i.strip() for i in open("positive-words.txt", "r").readlines()])
    negWords_list = set([i.strip() for i in open("negative-words.txt", "r", encoding="ISO-8859-1").readlines()])

    # count the number of positive and negative words present in each review text
    dict_ = defaultdict(list)
    for review_text in text:

        # remove punctuation symbols and spaces
        words = [n.lower().translate(str.maketrans('', '', string.punctuation)) for n in review_text.split(' ')]
        words = [i for i in words if i != ''] # remove spaces

        # count number of positive and negative words in each review
        pos_count, neg_count = 0, 0
        for word in words:
            if word in posWords_list:
                pos_count+=1
            elif word in negWords_list:
                neg_count+=1
            else:
                continue

        # save count and rate
        if len(words) == 0: # prevent division by zero
            dict_['pos_count'].append(-1)
            dict_['neg_count'].append(-1)
            dict_['pos_rate'].append(-1)
            dict_['neg_rate'].append(-1)
        else:
            dict_['pos_count'].append(pos_count)
            dict_['neg_count'].append(neg_count)
            dict_['pos_rate'].append(pos_count/len(words))
            dict_['neg_rate'].append(neg_count/len(words))

    return dict_
```



```
In [5]: def Get_NumWords(text):

        """
        This function counts the number of words in the text by:
            - words are converted to lower-case
            - blank spaces and punctuations are removed
            - stopwords are removed using NLTK library
        """

        # get stopwords from nltk library
        # nltk.download('stopwords')
        stop_words = set(stopwords.words('english'))

        # loop over each review
        list_ = []
        for review_text in text:

            # remove punctuation symbols and spaces
            words = [n.lower().translate(str.maketrans('', '', string.punctuation)) for n in review_text.split(' ')]
            words = [i for i in words if i != ''] # remove spaces

            # count number of words, excluding stopwords
            word_count = 0
            for word in words:
                if word in stop_words:
                    continue
                else:
                    word_count+=1

            # save counts
            list_.append(word_count)

        return list_
```

```
In [6]: def Get_category_numtrans(category_lists):

    """
    This function assigns an ID number to each unique list of categories,
    then those IDs are used to create a sum for the collection of lists
    found in each product Category column.
    """

    # get unique category lists
    lists = []
    for i in category_lists:
        for j in i:
            lists.append(tuple(j))

    # make dictionary with IDs
    categorylists_dict = {k:v for v,k in enumerate(set(lists))}

    # transform list occurrence to numbers using dictionary
    category_numtrans = []
    for i in category_lists:
        sum_ = 0
        for j in i:
            sum_ += categorylists_dict[tuple(j)]
        category_numtrans.append(sum_)

    # return transformation
    return category_numtrans
```

```
In [7]: def get_helpfulRate(col):

    """
    This function computes the ratio between number
    of helpfull votes and total votes.
    """

    allHelpful = []
    colHelpful = defaultdict(list)

    col_data = X_train_raw[col]
    allHelpful = y_train_raw

    for x,y in zip(col_data, allHelpful):
        colHelpful[x].append(y)

    averageRate = sum([x['nHelpful'] for x in allHelpful]) * 1.0 / sum([x['outOf'] for x in allHelpful])

    rate = {}
    for u in colHelpful:
        totalU = sum([x['outOf'] for x in colHelpful[u]])
        if totalU > 0:
            rate[u] = sum([x['nHelpful'] for x in colHelpful[u]]) * 1.0 / totalU
        else:
            rate[u] = averageRate

    return rate, averageRate
```

```
In [8]: def find_Helpful_rate(data, colname):

        """
        This function computes the helpful rate for a specified column
        For example: 'reviewerID' and 'itemID'.
        """

        rate_dict, avg_rate = get_helpfulRate(colname)

        ratehelpful = []
        for i in data:
            # use average for entries not present
            try:
                ratehelpful.append(rate_dict[i])
            except:
                ratehelpful.append(avg_rate)

        return ratehelpful
```

```
In [9]: def Get_punctuation_count(data):

        """
        This function counts the number of punctuation characters found in the text
        punctuation characters are taken from the 'string' library.
        """

        punct_set = set(string.punctuation)
        list_ = []
        for str_ in data:
            count_ = 0
            for c in str_:
                if c in punct_set:
                    count_ += 1
                else:
                    continue
            list_.append(count_)

        return list_
```

```
In [10]: def Get_numreviews_summarized(data):

    """
    This function finds the number of reviews encountered for the column specified:
    For example: 'itemID' and 'reviewerID'.
    """

    # initialize dict
    dict_ = {k:0 for k in set(data)}

    # count number
    for r in data:
        dict_[r]+=1

    # create list
    list_ = []
    for r in data:
        try:
            list_.append(dict_[r])
        except:
            list_.append(0)

    return list_
```

```
In [11]: def Get_product_ratingDeviation(df):

    """
    This function finds the difference in rating for each review
    as compared to the mean of all ratings for the item.
    """

    ProdRating_mean_dict = df[['itemID', 'rating']].groupby(['itemID']).mean().rating.to_dict()

    rating_deviation = []
    for r, p in zip(df.rating, df.itemID):
        try:
            rating_deviation.append(r - ProdRating_mean_dict[p])
        except:
            rating_deviation.append(-444)

    return rating_deviation
```

```
In [12]: def Get_count_firstLetterCapital(data):

    """
    This function counts the number of words in the text
    whose first letter is capital.
    """

    list_ = []
    for review in data:
        count_ = 0
        for word in review.split():
            if word[0].isupper():
                count_+=1
        list_.append(count_)

    return list_
```

```
In [13]: def Get_avg_word_length(data):

    """
    This function computes the average word length for the text,
    stop words are removed using the list from NLTK library.
    """

    # get stopwords from nltk library
    # nltk.download('stopwords')
    stop_words = set(stopwords.words('english'))

    # loop over each review
    list_ = []
    for review_text in data:

        # remove punctuation symbols and spaces
        words = [n.lower().translate(str.maketrans('', '', string.punctuation)) for n in review_text.split(' ')]
        words = [i for i in words if i != ''] # remove spaces

        # remove stopwords
        words = set(words) - stop_words

        # save average length
        try:
            list_.append(sum([len(w) for w in words])/len(words))
        except:
            list_.append(-1)

    return list_
```

```
In [14]: def Get_delta_sinceFirstReview(df):

    """
    This function finds the Unix time difference
    since the first review for each item.
    """

    # find time for first review in each product
    first_product_reviewtime = df[['itemID', 'unixReviewTime']].groupby(['itemID']).min().unixReviewTime.to_dict()

    # find delta for each review
    delta = []
    for i,t in zip(df.itemID, df.unixReviewTime):
        delta.append(t-first_product_reviewtime[i])

    return delta
```



```
In [15]: def Get_delta_sinceLastReview(df):

    """
    This function finds the Unix time difference
    since the last review for each item.
    """

    # create reduced df
    new_df = df[['itemID', 'unixReviewTime']].groupby(['itemID'])['unixReviewTime'].apply(list)

    # create dictionary with deltas
    delta_dict = defaultdict(dict)

    for idx, times_list in zip(new_df.index, new_df):

        times_list = sorted(times_list)
        times_list_deltas = np.append(np.array([0]), np.diff(times_list))

        for t, d in zip(times_list, times_list_deltas):
            delta_dict[idx][t] = d

    # generate list with deltas matching input dataframe
    list_ = []
    for u, t in zip(df.itemID, df.unixReviewTime):
        list_.append(delta_dict[u][t])

    return list_
```

```
In [16]: def Get_colsRatio(col1, col2):

    """
    This function finds the ratio element-wise for two columns.
    """

    list_ = []
    for i, j in zip(col1, col2):

        try:
            list_.append(i/j)
        except:
            list_.append(-1)

    return list_
```

```
In [17]: def Get_numCapitalwords(data):

    """
    This function finds the number of words
    in the text where all letters are capital.
    """

    # get stopwords from nltk library
    # nltk.download('stopwords')
    stop_words = set(stopwords.words('english'))

    # loop over each review
    list_ = []
    for review_text in data:

        # remove punctuation symbols and spaces
        words = [n.translate(str.maketrans('', '', string.punctuation)) for n in review_text.split(' ')]
        words = [i for i in words if i != ''] # remove spaces

        # remove stopwords
        words = set(words) - stop_words

        # append count of upper case words
        list_.append(len([i for i in words if i.isupper()]))

    return list_
```

```
In [18]: def Get_ExclQues_charCount(data):

    """
    This function counts the number of exclamation
    and question characters in the text.
    """

    # loop over each review
    list_ = []
    for review_text in data:
        # append count of question and exclamation characters
        list_.append(review_text.count('?') + review_text.count('!'))

    return list_
```

```
In [19]: def Get_features(df):

    """
    This function learns all the features from the raw data,
    each feature is added as a new column to the input dataframe.
    """

    # Modify -----> "categories"

    # get number of characters
    df.loc[:, 'categories_count'] = [len(i) for i in df['categories']]

    # generate numerical category by transforming combination of lists to numbers
    df.loc[:, 'category_numtrans'] = Get_category_numtrans(df.categories)

    # Modify -----> "itemID" and "reviewerID"

    # create dictionaries for itemID and reviewerID, convert from categorical to numeric
    items_dict = {k:v for v,k in enumerate(set(df.itemID))}
    reviewer_dict = {k:v for v,k in enumerate(set(df.reviewerID))}

    # change item and reviewer IDs to numeric
    df.loc[:, 'itemID'] = [items_dict[i] for i in df['itemID']]
    df.loc[:, 'reviewerID'] = [reviewer_dict[i] for i in df['reviewerID']]

    # add helpful rate for itemID and reviewerID
    df.loc[:, 'itemID_helpfulRate'] = find_Helpful_rate(df['itemID'], 'itemID')
    df.loc[:, 'reviewerID_helpfulRate'] = find_Helpful_rate(df['reviewerID'], 'reviewerID')

    # get number of reviews for each user
    df.loc[:, 'reviewerID_numReviews'] = Get_numreviews_summarized(df['reviewerID'])

    # get number of reviews for each product
    df.loc[:, 'itemID_numReviews'] = Get_numreviews_summarized(df['itemID'])

    # Modify -----> "reviewText"

    # get number of words, remove stopwords
    df.loc[:, 'reviewText_count_words'] = Get_NumWords(df['reviewText'])

    # get character count
    df.loc[:, 'reviewText_count_char'] = [len(i) for i in df['reviewText']]

    # get punctuation count
```

```

df.loc[:, 'reviewText_count_punctu'] = Get_punctuation_count(df['reviewText'])

# get number of words that start with a capital letter
df.loc[:, 'reviewText_count_firstCapital'] = Get_count_firstLetterCapital(df['reviewText'])

# get average word length
df.loc[:, 'reviewText_avgWordLength'] = Get_avg_word_length(df['reviewText'])

# get number of capital words
df.loc[:, 'reviewText_capitalwords'] = Get_numCapitalwords(df['reviewText'])

# get number of question and exclamation characters
df.loc[:, 'reviewText_ExclQue_countchar'] = Get_ExclQues_charCount(df['reviewText'])

# get ratio between punctuations with character numbers
df.loc[:, 'reviewText_PunctChar_ratio'] = Get_colsRatio(df['reviewText_count_punctu'], df['reviewText_count_char'])

# get positive and negative word rate
reviewText_PosNeg = GetPosNegWords_count_Rate(df.reviewText)
df.loc[:, 'reviewText_posWordCount'] = reviewText_PosNeg['pos_count']
df.loc[:, 'reviewText_negWordCount'] = reviewText_PosNeg['neg_count']
df.loc[:, 'reviewText_posWordRate'] = reviewText_PosNeg['pos_rate']
df.loc[:, 'reviewText_negWordRate'] = reviewText_PosNeg['neg_rate']

# Modify -----> "summary"

# get number of words, remove stopwords
df.loc[:, 'summary_count_words'] = Get_NumWords(df['summary'])

# get character count
df.loc[:, 'summary_count_char'] = [len(i) for i in df['summary']]

# get punctuation count
df.loc[:, 'summary_count_punctu'] = Get_punctuation_count(df['summary'])

# get number of words that start with a capital letter
df.loc[:, 'summary_count_firstCapital'] = Get_count_firstLetterCapital(df['summary'])

# get average word length
df.loc[:, 'summary_avgWordLength'] = Get_avg_word_length(df['summary'])

# get number of capital words
df.loc[:, 'summary_capitalwords'] = Get_numCapitalwords(df['summary'])

# get number of question and exclamation characters

```

```

df.loc[:, 'summary_ExclQue_countchar'] = Get_ExclQues_charCount(df['summary'])

# get ratio between punctuations with character numbers
df.loc[:, 'summary_PunctChar_ratio'] = Get_colsRatio(df['summary_count_punctu'], df['summary_count_char'])

# get positive and negative word rate
summary_PosNeg = GetPosNegWords_count_Rate(df.summary)
df.loc[:, 'summary_posWordCount'] = summary_PosNeg['pos_count']
df.loc[:, 'summary_negWordCount'] = summary_PosNeg['neg_count']
df.loc[:, 'summary_posWordRate'] = summary_PosNeg['pos_rate']
df.loc[:, 'summary_negWordRate'] = summary_PosNeg['neg_rate']

# Modify -----> "helpful"

# parse helpful votes
df.loc[:, 'outOf_feature'] = [i['outOf'] for i in df['helpful']]

# Modify -----> "price"

# change NA values to -999
df.loc[df.price.isna(), 'price'] = -999

# Modify -----> "unixReviewTime"

# find time since first review
df.loc[:, 'unixReviewTime_delta_firstreview'] = Get_delta_sinceFirstReview(df)

# find time since last review for same product
df.loc[:, 'unixReviewTime_delta_lastreview'] = Get_delta_sinceLastReview(df)

# Modify -----> "reviewText"

# get rating deviation from the mean
df.loc[:, 'rating_deviation'] = Get_product_ratingDeviation(df)

# Add -----> New Columns

# votes over time
df.loc[:, 'votes_time'] = df.outOf_feature/df.unixReviewTime

# ratio of summary to reviewText for characters and words
df.loc[:, 'summary_reviewText_charRatio'] = Get_colsRatio(df.summary_count_char, df.reviewText_count_char)
df.loc[:, 'summary_reviewText_wordsRatio'] = Get_colsRatio(df.summary_count_words, df.reviewText_count_words)

# define columns to keep

```

```

cols = ['categoryID', 'categories_count', 'category_numtrans', 'summary_count_words', 'itemID', 'reviewerID', 'rating',
        'reviewText_count_words', 'itemID_helpfulRate', 'reviewerID_helpfulRate',
        'reviewerID_numReviews', 'rating_deviation',
        'outOf_feature', 'unixReviewTime', 'price', 'reviewText_posWordCount', 'reviewText_negWordCount',
        'reviewText_posWordRate', 'reviewText_negWordRate', 'summary_count_char', 'reviewText_count_char',
        'reviewText_count_punctu', 'summary_count_punctu', 'summary_posWordCount',
        'summary_negWordCount', 'summary_posWordRate', 'summary_negWordRate',
        'reviewText_count_firstCapital', 'summary_count_firstCapital',
        'reviewText_avgWordLength', 'summary_avgWordLength', 'unixReviewTime_delta_firstreview',
        'votes_time', 'summary_reviewText_charRatio', 'summary_reviewText_wordsRatio', 'itemID_numReviews',
        'reviewText_capitalwords', 'summary_capitalwords',
        'reviewText_ExclQue_countchar', 'summary_ExclQue_countchar',
        'reviewText_PunctChar_ratio', 'summary_PunctChar_ratio', 'unixReviewTime_delta_lastreview']

# return features
return df[cols]

```

```

In [20]: def Get_labels_ratio(df):

        """
        This function finds the ratio between helpful votes
        with total votes from raw data. This ratio is the
        regression label used to train the model.
        """

        return [i['nHelpful']/i['outOf'] if i['outOf']!=0 else 0 for i in df]

```

```
In [21]: def save_predictions(pred):

    """
    This function reads a formatted file and writes
    the predictions found for the test data. The file
    is used as input to the DSE 220 Kaggle competition.
    """

    predictions = open("predictions_Helpful.txt", 'w')

    count = 0
    for l in open("pairs_Helpful.txt"):

        if l.startswith("userID"):
            #header
            predictions.write(l)
            continue

        u,i,outOf = l.strip().split('-')

        predictions.write(u + '-' + i + '-' + str(outOf) + ',' + str(pred[count]) + '\n')

        count+=1

    predictions.close()
```

```
In [22]: def perfmetrics_RFmodel(max_depth, n_estimators, X_train_filt, y_train_filt, X_train, y_train, X_test, y_test):
```

```
    """
```

```
    This function applies the random forest model,  
    the inputs and outputs are listed below:
```

```
    Inputs:
```

```
        Random Forest Hyperparameters
```

- max\_depth: max depth of the tree
- n\_estimators: number of trees in the forest

```
        Trainning Data
```

- X\_train\_filt: train the model with train data where number of votes (labels) are > 0
- y\_train\_filt train the model with train data where number of votes (labels) are > 0
- X\_train: data includes all labels
- y\_train: data includes all labels

```
        Testing Data (or Validation)
```

- X\_test: data includes all labels
- y\_test: data includes all labels

```
    Outputs:
```

- train\_accuracy
- test\_accuracy
- train\_mae (Mean Absolute Error)
- test\_mae
- train\_mse (Mean Squared Error)
- test\_mse
- y\_pred\_test (or validation predictions)

```
    """
```

```
    # define model
```

```
    regr = RandomForestRegressor(max_depth=max_depth, n_estimators=n_estimators, n_jobs=-1)
```

```
    # evaluate, with filtered data
```

```
    regr.fit(X_train_filt, y_train_filt)
```

```
    # predict test and train, multiply prediction by the number of votes
```

```
    y_pred_train = regr.predict(X_train)*np.array(X_train.outOf_feature)
```

```
    y_pred_test = regr.predict(X_test)*np.array(X_test.outOf_feature)
```

```
    # round to the nearest integer, labels are integers
```

```
    y_pred_train = [int(round(i)) for i in y_pred_train]
```

```
    y_pred_test = [int(round(i)) for i in y_pred_test]
```

```
    # calculate accuracy, mean absolute error, and mean squared error
```



```
train_accuracy = accuracy_score(y_train, y_pred_train)
train_mae = mean_absolute_error(y_train, y_pred_train)
train_mse = mean_squared_error(y_train, y_pred_train)

if y_test != None: # case when test labels are known
    test_accuracy = accuracy_score(y_test, y_pred_test)
    test_mae = mean_absolute_error(y_test, y_pred_test)
    test_mse = mean_squared_error(y_test, y_pred_test)
else: # case when test labels are predicted
    return (train_accuracy, train_mae, train_mse, y_pred_test)

# return all performance metrics
return (train_accuracy, test_accuracy, train_mae, test_mae, train_mse, test_mse, y_pred_test)
```

## Main Section

### Read Data

```
In [23]: %%time

# read train
train = getDF('train.json.gz')

# define features and labels for train data
train_features = train
train_labels = train['helpful']

# read test
test = getDF('test_Helpful.json.gz')

CPU times: user 27.3 s, sys: 1.47 s, total: 28.8 s
Wall time: 31.1 s
```

```
In [24]: # split train into train and validation
X_train_raw, X_val_raw, y_train_raw, y_val_raw = train_test_split(train_features, train_labels, test_size=0.20)
```

Table 1. Training data example (raw).

```
In [25]: X_train_raw.head()
```

Out[25]:

	categoryID	categories	itemID	reviewerID	rating	reviewText	reviewHash	reviewTime	summary	unixReviewTime	helpful	price
44011	0	[[Clothing, Shoes & Jewelry, D, Dreams], [Clot...	I322658212	U656051947	5.0	These pajamas are a perfect weight for me. I'm...	R196940592	03 25, 2014	Nice weight & comfy.	1395705600	{'outOf': 0, 'nHelpful': 0}	NaN
53593	0	[[Clothing, Shoes & Jewelry, Shoes & Accessori...	I913405870	U604395367	3.0	I returned these because I thought they looked...	R612741609	09 16, 2013	Nice quality	1379289600	{'outOf': 0, 'nHelpful': 0}	NaN
141214	0	[[Clothing, Shoes & Jewelry, Women, Shoes, San...	I036574902	U886882212	3.0	Saddle is very pretty, and I think it will be ...	R077184873	05 5, 2014	Very pretty	1399248000	{'outOf': 0, 'nHelpful': 0}	NaN
179392	0	[[Clothing, Shoes & Jewelry, Women], [Clothing...	I134433773	U154865904	5.0	This ring is truly stunning and dramatic while...	R611116635	03 16, 2014	Stainless Steel Band With CZ	1394928000	{'outOf': 0, 'nHelpful': 0}	5.94
32982	1	[[Clothing, Shoes & Jewelry, Men, Accessories,...	I532478566	U578222660	4.0	At first the tension on this was too tight, bu...	R286930585	05 16, 2014	No nonsense, just how I like it.	1400198400	{'outOf': 0, 'nHelpful': 0}	7.99

Data Preparation - Train and Validation

```
In [26]: %%time
```

```
# learn features
X_train = Get_features(X_train_raw)
X_val = Get_features(X_val_raw)

# create labels
y_train = [i['nHelpful'] for i in y_train_raw]
y_val = [i['nHelpful'] for i in y_val_raw]
```

CPU times: user 2min 39s, sys: 2.78 s, total: 2min 42s  
Wall time: 2min 43s

```
In [27]: %%time

# model is only trained for reviews where number of votes is > 0

# find indices
idx = np.array([i['outOf'] for i in X_train_raw.helpful]) > 0

# filter new train data and learn features on filtered data
X_train_filt = Get_features(X_train_raw.loc[idx,:])
y_train_ratio_filt = Get_labels_ratio(list(compress(y_train_raw, idx)))
```

CPU times: user 52 s, sys: 302 ms, total: 52.3 s

Wall time: 52.6 s

## Feature Selection - Keep Features with Better Performance than Baseline

```

In [28]: %%time

# define hyperparameters to test
max_depth = 10
n_estimators = 100

# define baseline and other features
baseline_features = ['itemID', 'reviewerID', 'outOf_feature']
features = [i for i in X_train.columns if i not in baseline_features]

# run model with baseline features + one feature, see MAE if improves
metrics = []
for i in range(len(features)+1):

    # define features to test
    if i == 0: # baseline
        curr_features = baseline_features
        feature_name = 'baseline'
    else: # other features
        curr_features = baseline_features + [features[i-1]]
        feature_name = features[i-1]

    # get metrics (train_accuracy, test_accuracy, train_mae, test_mae, train_mse, test_mse, y_pred_test)
    metrics_tuple = perfmetrics_RFmodel(max_depth, n_estimators, X_train_filt[curr_features],
                                         y_train_ratio_filt, X_train[curr_features], y_train,
                                         X_val[curr_features], y_val)

    # save metrics
    metrics.append((feature_name,) + metrics_tuple[:-1])

```

CPU times: user 12min 7s, sys: 7.57 s, total: 12min 15s  
Wall time: 3min 43s

**Table 2. Features tested one-by-one along with baseline, those with higher than baseline performance (lower MAE) are selected.**

```
In [29]: # create dataframe
df_metrics = pd.DataFrame(metrics, columns=['features', 'train_accur', 'val_accur', 'train_mae', 'val_mae',
                                             'train_mse', 'val_mse'])

# sort by MAE in validation, descending
df_metrics = df_metrics.sort_values(by='val_mae')

df_metrics
```

Out[29]:

	features	train_accur	val_accur	train_mae	val_mae	train_mse	val_mse
10	rating_deviation	0.857513	0.863350	0.183419	0.179500	0.753856	0.742850
21	summary_posWordCount	0.858394	0.859800	0.182506	0.183150	0.615256	0.812350
23	summary_posWordRate	0.856988	0.858400	0.183706	0.183825	0.643056	0.802025
11	unixReviewTime	0.855681	0.857625	0.183688	0.184100	0.582037	0.828750
15	reviewText_posWordRate	0.856012	0.858625	0.181719	0.184650	0.514431	0.914550
5	rating	0.859163	0.864450	0.178381	0.184875	0.482044	1.891625
24	summary_negWordRate	0.855675	0.857600	0.184012	0.184975	0.616337	0.858525
13	reviewText_posWordCount	0.855444	0.856925	0.184200	0.185400	0.524800	0.831150
22	summary_negWordCount	0.853819	0.856400	0.186306	0.185575	0.627181	0.846875
38	reviewText_PunctChar_ratio	0.854006	0.855700	0.186356	0.185800	0.554556	0.836600
18	reviewText_count_char	0.857587	0.857500	0.180156	0.185800	0.554769	0.914700
12	price	0.854738	0.856050	0.184037	0.185850	0.553400	0.857850
36	reviewText_ExclQue_countchar	0.853450	0.854650	0.187562	0.186000	0.685662	0.841750
31	summary_reviewText_charRatio	0.854962	0.855250	0.184113	0.186175	0.617812	0.879275
6	reviewText_count_words	0.856363	0.856500	0.182669	0.186300	0.607556	0.888850
37	summary_ExclQue_countchar	0.853762	0.854300	0.185938	0.186375	0.623075	0.849875
29	unixReviewTime_delta_firstreview	0.854838	0.854300	0.184275	0.186425	0.657975	0.854025
1	categoryID	0.847069	0.856250	0.193944	0.186500	0.687056	0.848750
25	reviewText_count_firstCapital	0.855750	0.854925	0.184306	0.186550	0.617419	0.846350
26	summary_count_firstCapital	0.855481	0.855225	0.187044	0.186625	0.674181	0.845625

	features	train_accur	val_accur	train_mae	val_mae	train_mse	val_mse
35	summary_capitalwords	0.854475	0.853825	0.186550	0.186725	0.665575	0.836625
39	summary_PunctChar_ratio	0.854675	0.854650	0.184806	0.186825	0.661294	0.849475
20	summary_count_punctu	0.853744	0.854475	0.187250	0.186900	0.677963	0.847800
40	unixReviewTime_delta_lastreview	0.854044	0.854100	0.186406	0.186900	0.647744	0.851700
28	summary_avgWordLength	0.855537	0.854900	0.183356	0.186925	0.608669	0.877225
19	reviewText_count_punctu	0.855206	0.853725	0.184331	0.187200	0.549556	0.739850
32	summary_reviewText_wordsRatio	0.854050	0.854550	0.185087	0.187500	0.623687	0.881500
14	reviewText_negWordCount	0.854550	0.853525	0.186031	0.187500	0.654481	0.846400
33	itemID_numReviews	0.851650	0.852875	0.201437	0.187550	0.829775	0.857650
3	category_numtrans	0.853400	0.853425	0.187638	0.187575	0.558163	0.858025
30	votes_time	0.853525	0.852875	0.187888	0.187750	0.688375	0.866300
34	reviewText_capitalwords	0.853900	0.852875	0.185938	0.187875	0.655900	0.853925
8	reviewerID_helpfulRate	0.857487	0.853150	0.182163	0.187925	0.671488	0.857475
0	baseline	0.853556	0.852500	0.186312	0.188325	0.677188	0.865975
4	summary_count_words	0.854725	0.853525	0.186587	0.188350	0.650500	0.866250
7	itemID_helpfulRate	0.856669	0.852975	0.184888	0.188400	0.697700	0.860200
17	summary_count_char	0.854369	0.853175	0.186069	0.188425	0.657131	0.869125
2	categories_count	0.853712	0.852300	0.188206	0.188775	0.575481	0.860575
16	reviewText_negWordRate	0.855213	0.856225	0.183631	0.189450	0.501481	1.175300
9	reviewerID_numReviews	0.853731	0.851750	0.187300	0.189550	0.643663	0.859650
27	reviewText_avgWordLength	0.854031	0.853475	0.185463	0.190525	0.592287	1.120425

```
In [30]: # keep features with better performance than baseline
baseline_valMAE = df_metrics.val_mae[df_metrics.features == 'baseline'].values[0]

features_aboveBaseline = list(df_metrics.features[df_metrics.val_mae < baseline_valMAE])
all_features = baseline_features + features_aboveBaseline

print(all_features)

['itemID', 'reviewerID', 'outOf_feature', 'rating_deviation', 'summary_posWordCount', 'summary_posWordRate', 'unixReviewTime', 'reviewText_posWordRate', 'rating', 'summary_negWordRate', 'reviewText_posWordCount', 'summary_negWordCount', 'reviewText_PunctChar_ratio', 'reviewText_count_char', 'price', 'reviewText_ExclQue_countchar', 'summary_reviewText_charRatio', 'reviewText_count_words', 'summary_ExclQue_countchar', 'unixReviewTime_delta_firstreview', 'categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

## Feature Selection - Ablation Study with Features Elimination (Ordered by MAE descending)

In [31]: %%time

```
# define hyperparameters to test
max_depth = 10
n_estimators = 100

# save list of features eliminated and the calculated MAE
features_list_mae_mse = []

# eliminate each feature, one-by-one, starting with the highest MAE feature to the lowest on the validation dataset
for f in range(len(all_features)-2): # features list is sorted with MAE ascending

    # eliminate starting from the end
    temp_featurelist = all_features[:len(all_features)-f]

    # get feature name
    if f == 0:
        f_eliminated = 'None'
    else:
        f_eliminated = all_features[-f:]

    # get metrics (train_accuracy, test_accuracy, train_mae, test_mae, train_mse, test_mse, y_pred_test)
    metrics_tuple = perfmetrics_RFmodel(max_depth, n_estimators, X_train_filt[temp_featurelist],
                                         y_train_ratio_filt, X_train[temp_featurelist],
                                         y_train, X_val[temp_featurelist], y_val)

    # print metrics
    print('Features Eliminated: {}'.format(f_eliminated))
    print('\tTrain Accuracy: {}'.format(metrics_tuple[0]))
    print('\tValidation Accuracy: {}'.format(metrics_tuple[1]))
    print('\tTrain MAE: {}'.format(metrics_tuple[2]))
    print('\tValidation MAE: {}'.format(metrics_tuple[3]))
    print('\tTrain MSE: {}'.format(metrics_tuple[4]))
    print('\tValidation MSE: {}'.format(metrics_tuple[5]))
    print('\n')

    # save mae and eliminated features lists
    features_list_mae_mse.append((metrics_tuple[3], metrics_tuple[5], f_eliminated))
```

```
Features Eliminated: None
    Train Accuracy: 0.8679625
    Validation Accuracy: 0.86735
    Train MAE: 0.1681
    Validation MAE: 0.1741
```



Train MSE: 0.5105125  
Validation MSE: 0.769

Features Eliminated: ['reviewerID\_helpfulRate']

Train Accuracy: 0.86815  
Validation Accuracy: 0.867725  
Train MAE: 0.1678125  
Validation MAE: 0.173675  
Train MSE: 0.5429125  
Validation MSE: 0.772075

Features Eliminated: ['reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8676375  
Validation Accuracy: 0.8675  
Train MAE: 0.1694  
Validation MAE: 0.17405  
Train MSE: 0.5346125  
Validation MSE: 0.71645

Features Eliminated: ['votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.86820625  
Validation Accuracy: 0.86715  
Train MAE: 0.1682625  
Validation MAE: 0.17425  
Train MSE: 0.54125  
Validation MSE: 0.72675

Features Eliminated: ['category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8680375  
Validation Accuracy: 0.86735  
Train MAE: 0.168725  
Validation MAE: 0.1749  
Train MSE: 0.5389875  
Validation MSE: 0.80905

Features Eliminated: ['itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8701875  
Validation Accuracy: 0.8668  
Train MAE: 0.162575  
Validation MAE: 0.17375

Train MSE: 0.48335  
Validation MSE: 0.7025

Features Eliminated: ['reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.87045  
Validation Accuracy: 0.86705  
Train MAE: 0.16189375  
Validation MAE: 0.173525  
Train MSE: 0.51855625  
Validation MSE: 0.708675

Features Eliminated: ['summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.87003125  
Validation Accuracy: 0.866925  
Train MAE: 0.16253125  
Validation MAE: 0.173325  
Train MSE: 0.50564375  
Validation MSE: 0.708675

Features Eliminated: ['reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.870525  
Validation Accuracy: 0.8671  
Train MAE: 0.1616125  
Validation MAE: 0.173825  
Train MSE: 0.479925  
Validation MSE: 0.715575

Features Eliminated: ['summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.869575  
Validation Accuracy: 0.866725  
Train MAE: 0.16244375  
Validation MAE: 0.174  
Train MSE: 0.50869375  
Validation MSE: 0.73725

Features Eliminated: ['unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRati

```
o', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.87056875
Validation Accuracy: 0.867375
Train MAE: 0.16161875
Validation MAE: 0.17375
Train MSE: 0.48356875
Validation MSE: 0.7448
```

```
Features Eliminated: ['summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.870425
Validation Accuracy: 0.867075
Train MAE: 0.16193125
Validation MAE: 0.1733
Train MSE: 0.48360625
Validation MSE: 0.7058
```

```
Features Eliminated: ['summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.8704125
Validation Accuracy: 0.867275
Train MAE: 0.1619375
Validation MAE: 0.1736
Train MSE: 0.5141
Validation MSE: 0.735
```

```
Features Eliminated: ['summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.869875
Validation Accuracy: 0.867325
Train MAE: 0.16215
Validation MAE: 0.17355
Train MSE: 0.4518375
Validation MSE: 0.7451
```

```
Features Eliminated: ['summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'i
```

```
temID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.8698
```

```
Validation Accuracy: 0.867225
```

```
Train MAE: 0.1622875
```

```
Validation MAE: 0.173975
```

```
Train MSE: 0.4795625
```

```
Validation MSE: 0.744025
```

```
Features Eliminated: ['reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.86995
```

```
Validation Accuracy: 0.8667
```

```
Train MAE: 0.1619125
```

```
Validation MAE: 0.174425
```

```
Train MSE: 0.47415
```

```
Validation MSE: 0.744225
```

```
Features Eliminated: ['categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.86934375
```

```
Validation Accuracy: 0.8664
```

```
Train MAE: 0.1634125
```

```
Validation MAE: 0.176625
```

```
Train MSE: 0.554575
```

```
Validation MSE: 0.921025
```

```
Features Eliminated: ['unixReviewTime_delta_firstreview', 'categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.869575
```

```
Validation Accuracy: 0.867475
```

```
Train MAE: 0.16406875
```

```
Validation MAE: 0.175425
```

```
Train MSE: 0.57839375
```

```
Validation MSE: 0.907625
```

Features Eliminated: ['summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.86935625  
Validation Accuracy: 0.8679  
Train MAE: 0.16398125  
Validation MAE: 0.174725  
Train MSE: 0.54790625  
Validation MSE: 0.894625

Features Eliminated: ['reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.86899375  
Validation Accuracy: 0.867075  
Train MAE: 0.16439375  
Validation MAE: 0.176375  
Train MSE: 0.52345625  
Validation MSE: 0.984725

Features Eliminated: ['summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.868775  
Validation Accuracy: 0.867775  
Train MAE: 0.16463125  
Validation MAE: 0.17545  
Train MSE: 0.53184375  
Validation MSE: 0.96365

Features Eliminated: ['reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.86899375  
Validation Accuracy: 0.8674  
Train MAE: 0.1639875

Validation MAE: 0.17685  
Train MSE: 0.512375  
Validation MSE: 1.0431

Features Eliminated: ['price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8684375  
Validation Accuracy: 0.867325  
Train MAE: 0.16644375  
Validation MAE: 0.1771  
Train MSE: 0.56099375  
Validation MSE: 0.9935

Features Eliminated: ['reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.86693125  
Validation Accuracy: 0.8672  
Train MAE: 0.16788125  
Validation MAE: 0.178075  
Train MSE: 0.54553125  
Validation MSE: 1.050725

Features Eliminated: ['reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8660875  
Validation Accuracy: 0.8668  
Train MAE: 0.16935  
Validation MAE: 0.178025  
Train MSE: 0.5467  
Validation MSE: 1.176475

Features Eliminated: ['summary\_negWordCount', 'reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countcha

```
r', 'summary_reviewText_charRatio', 'reviewText_count_words', 'summary_ExclQue_countchar', 'unixReviewTime_delta_firstreview', 'categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.8662375  
Validation Accuracy: 0.8667  
Train MAE: 0.168975  
Validation MAE: 0.17905  
Train MSE: 0.52705  
Validation MSE: 1.25515
```

```
Features Eliminated: ['reviewText_posWordCount', 'summary_negWordCount', 'reviewText_PunctChar_ratio', 'reviewText_count_char', 'price', 'reviewText_ExclQue_countchar', 'summary_reviewText_charRatio', 'reviewText_count_words', 'summary_ExclQue_countchar', 'unixReviewTime_delta_firstreview', 'categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.866  
Validation Accuracy: 0.867  
Train MAE: 0.17045  
Validation MAE: 0.178475  
Train MSE: 0.5603875  
Validation MSE: 1.141925
```

```
Features Eliminated: ['summary_negWordRate', 'reviewText_posWordCount', 'summary_negWordCount', 'reviewText_PunctChar_ratio', 'reviewText_count_char', 'price', 'reviewText_ExclQue_countchar', 'summary_reviewText_charRatio', 'reviewText_count_words', 'summary_ExclQue_countchar', 'unixReviewTime_delta_firstreview', 'categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

```
Train Accuracy: 0.86541875  
Validation Accuracy: 0.8664  
Train MAE: 0.1709375  
Validation MAE: 0.1783  
Train MSE: 0.5847125  
Validation MSE: 1.08385
```

```
Features Eliminated: ['rating', 'summary_negWordRate', 'reviewText_posWordCount', 'summary_negWordCount', 'reviewText_PunctChar_ratio', 'reviewText_count_char', 'price', 'reviewText_ExclQue_countchar', 'summary_reviewText_charRatio', 'reviewText_count_words', 'summary_ExclQue_countchar', 'unixReviewTime_delta_firstreview', 'categoryID', 'reviewText_count_firstCapital', 'summary_count_firstCapital', 'summary_capitalwords', 'summary_PunctChar_ratio', 'summary_count_punctu', 'unixReviewTime_delta_lastreview', 'summary_avgWordLength', 'reviewText_count_punctu', 'summary_reviewText_wordsRatio', 'reviewText_negWordCount', 'itemID_numReviews', 'category_numtrans', 'votes_time', 'reviewText_capitalwords', 'reviewerID_helpfulRate']
```

Train Accuracy: 0.86361875  
Validation Accuracy: 0.864475  
Train MAE: 0.1736625  
Validation MAE: 0.1783  
Train MSE: 0.729225  
Validation MSE: 0.97285

Features Eliminated: ['reviewText\_posWordRate', 'rating', 'summary\_negWordRate', 'reviewText\_posWordCount', 'summary\_negWordCount', 'reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8612375  
Validation Accuracy: 0.863875  
Train MAE: 0.17938125  
Validation MAE: 0.178175  
Train MSE: 0.77003125  
Validation MSE: 0.830275

Features Eliminated: ['unixReviewTime', 'reviewText\_posWordRate', 'rating', 'summary\_negWordRate', 'reviewText\_posWordCount', 'summary\_negWordCount', 'reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.85985625  
Validation Accuracy: 0.863275  
Train MAE: 0.1816125  
Validation MAE: 0.1796  
Train MSE: 0.7793125  
Validation MSE: 0.9308

Features Eliminated: ['summary\_posWordRate', 'unixReviewTime', 'reviewText\_posWordRate', 'rating', 'summary\_negWordRate', 'reviewText\_posWordCount', 'summary\_negWordCount', 'reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']

Train Accuracy: 0.8595125



Validation Accuracy: 0.864  
Train MAE: 0.1830875  
Validation MAE: 0.1788  
Train MSE: 0.8068875  
Validation MSE: 0.81785

Features Eliminated: ['summary\_posWordCount', 'summary\_posWordRate', 'unixReviewTime', 'reviewText\_posWordRate', 'rating', 'summary\_negWordRate', 'reviewText\_posWordCount', 'summary\_negWordCount', 'reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']  
Train Accuracy: 0.858275  
Validation Accuracy: 0.86325  
Train MAE: 0.1827875  
Validation MAE: 0.18015  
Train MSE: 0.77225  
Validation MSE: 0.7815

Features Eliminated: ['rating\_deviation', 'summary\_posWordCount', 'summary\_posWordRate', 'unixReviewTime', 'reviewText\_posWordRate', 'rating', 'summary\_negWordRate', 'reviewText\_posWordCount', 'summary\_negWordCount', 'reviewText\_PunctChar\_ratio', 'reviewText\_count\_char', 'price', 'reviewText\_ExclQue\_countchar', 'summary\_reviewText\_charRatio', 'reviewText\_count\_words', 'summary\_ExclQue\_countchar', 'unixReviewTime\_delta\_firstreview', 'categoryID', 'reviewText\_count\_firstCapital', 'summary\_count\_firstCapital', 'summary\_capitalwords', 'summary\_PunctChar\_ratio', 'summary\_count\_punctu', 'unixReviewTime\_delta\_lastreview', 'summary\_avgWordLength', 'reviewText\_count\_punctu', 'summary\_reviewText\_wordsRatio', 'reviewText\_negWordCount', 'itemID\_numReviews', 'category\_numtrans', 'votes\_time', 'reviewText\_capitalwords', 'reviewerID\_helpfulRate']  
Train Accuracy: 0.85410625  
Validation Accuracy: 0.852775  
Train MAE: 0.1873  
Validation MAE: 0.18855  
Train MSE: 0.6837875  
Validation MSE: 0.8667

CPU times: user 30min 25s, sys: 14.3 s, total: 30min 39s  
Wall time: 9min 39s

Figure 1. Ablation study results, MAE and MSE.

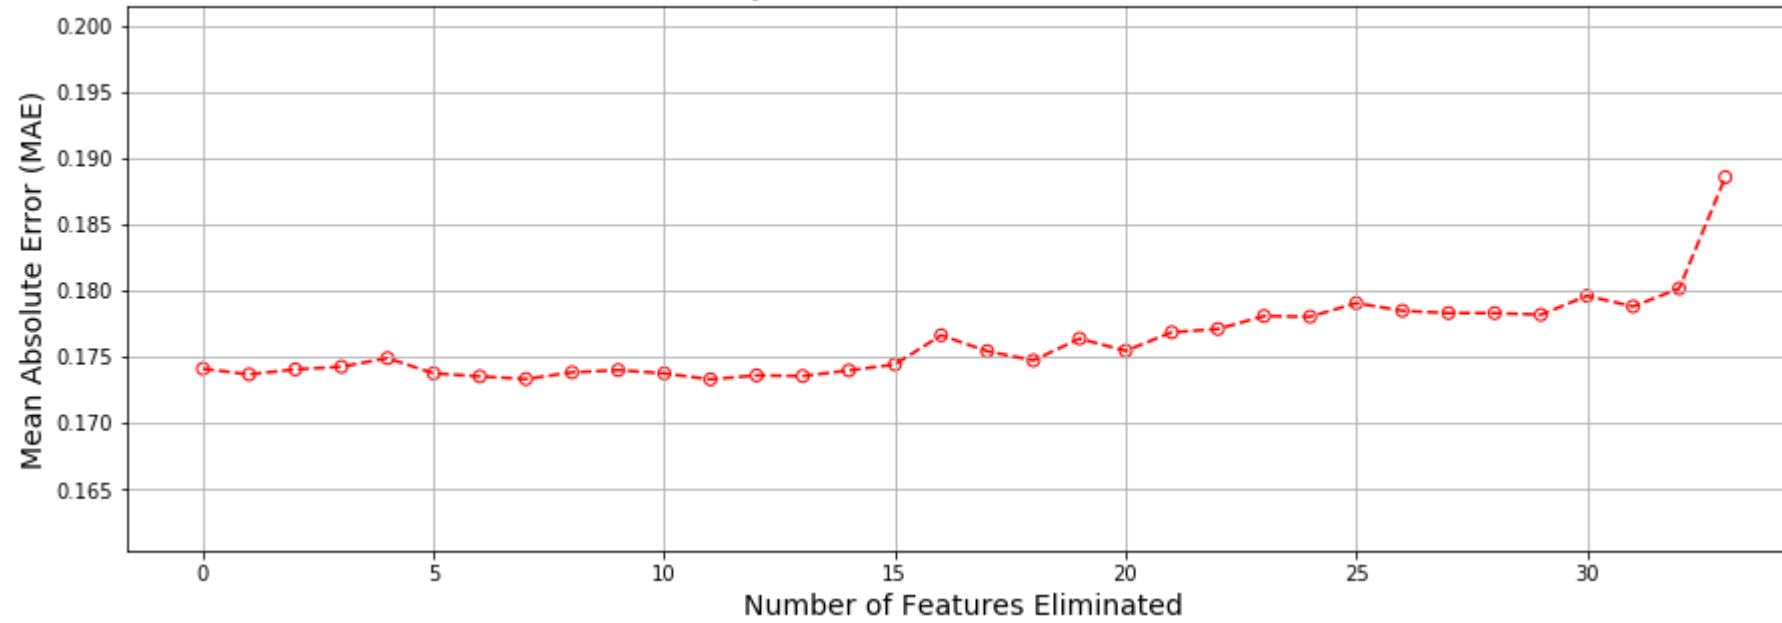
```
In [32]: # plot metrics vs. number of features removed

# get number of features and mae values
feature_counts = [len(i[2]) if i[2] != 'None' else 0 for i in features_list_mae_mse]
mae_values = [i[0] for i in features_list_mae_mse]
mse_values = [i[1] for i in features_list_mae_mse]


# plot MAE
plt.figure(figsize=(15,5)) # define figure
plt.scatter(feature_counts, mae_values, facecolors='none', edgecolors='r')
plt.plot(feature_counts, mae_values, '--', color='r')
plt.title('Ablation Study Results - One Feature Removed at a Time', fontsize=14)
plt.xlabel('Number of Features Eliminated', fontsize=14)
plt.ylabel('Mean Absolute Error (MAE)', fontsize=14)
plt.grid()
plt.show()

# plot MSE
plt.figure(figsize=(15,5)) # define figure
plt.scatter(feature_counts, mse_values, facecolors='none', edgecolors='b')
plt.plot(feature_counts, mse_values, '--', color='b')
plt.title('Ablation Study Results - One Feature Removed at a Time', fontsize=14)
plt.xlabel('Number of Features Eliminated', fontsize=14)
plt.ylabel('Mean Squared Error (MSE)', fontsize=14)
plt.grid()
plt.show()
```

Ablation Study Results - One Feature Removed at a Time



Ablation Study Results - One Feature Removed at a Time



```
In [33]: # define features to keep based on ablation study results, criteria = lowest MAE
features_tokeep = set(all_features) - set(sorted(features_list_mae_mse, key= lambda x: x[0])[0][2])

print('A total of {} features were eliminated and not included downstream'
      .format(len(all_features)-len(features_tokeep)))
print('\n{} new features were selected, they are listed below:'.format(len(features_tokeep)-3))
print(features_tokeep)
```

A total of 11 features were eliminated and not included downstream

22 new features were selected, they are listed below:

```
{'itemID', 'reviewText_count_char', 'summary_negWordCount', 'summary_PunctChar_ratio', 'summary_negWordRate', 'price', 'reviewerID', 'summary_capitalwords', 'reviewText_posWordRate', 'reviewText_ExclQue_countchar', 'outOf_feature', 'unixReviewTime', 'rating_deviation', 'summary_posWordCount', 'reviewText_PunctChar_ratio', 'summary_posWordRate', 'summary_ExclQue_countchar', 'categoryID', 'reviewText_posWordCount', 'unixReviewTime_delta_firstreview', 'summary_count_firstCapital', 'summary_reviewText_charRatio', 'rating', 'reviewText_count_words', 'reviewText_count_firstCapital'}
```

## Hyperparameter Tuning - Training and Validation Datasets

In [34]: %%time

```
# run test and validation dataset to tune parameters: max_depth, n_estimators
hyper_parameter_list = []
y_pred_val_list = []
for max_depth in [10,12]:
    for n_estimators in [100, 200, 400, 600]:

        # get metrics (train_accuracy, test_accuracy, train_mae, test_mae, train_mse, test_mse, y_pred_test)
        metrics_tuple = perfmetrics_RFmodel(max_depth, n_estimators, X_train_filt[features_tokeep],
                                             y_train_ratio_filt, X_train[features_tokeep],
                                             y_train, X_val[features_tokeep], y_val)

        # print metrics
        print('max_depth: {}, n_estimators:{}'.format(max_depth, n_estimators))
        print('\tTrain Accuracy: {}'.format(metrics_tuple[0]))
        print('\tValidation Accuracy: {}'.format(metrics_tuple[1]))
        print('\tTrain MAE: {}'.format(metrics_tuple[2]))
        print('\tValidation MAE: {}'.format(metrics_tuple[3]))
        print('\tTrain MSE: {}'.format(metrics_tuple[4]))
        print('\tValidation MSE: {}'.format(metrics_tuple[5]))
        print('\n')

        # save hyperparameters, mae, and predictions
        hyper_parameter_list.append((max_depth, n_estimators, metrics_tuple[3], metrics_tuple[6]))
```

```
max_depth: 10, n_estimators:100
    Train Accuracy: 0.87045625
    Validation Accuracy: 0.86675
    Train MAE: 0.1618625
    Validation MAE: 0.172775
    Train MSE: 0.5269875
    Validation MSE: 0.668575
```

```
max_depth: 10, n_estimators:200
    Train Accuracy: 0.87031875
    Validation Accuracy: 0.86695
    Train MAE: 0.1617375
    Validation MAE: 0.1735
    Train MSE: 0.492075
    Validation MSE: 0.7034
```

```
max_depth: 10, n_estimators:400
```

Train Accuracy: 0.87043125  
Validation Accuracy: 0.867  
Train MAE: 0.1617375  
Validation MAE: 0.1734  
Train MSE: 0.4919125  
Validation MSE: 0.70715

max\_depth: 10, n\_estimators:600  
Train Accuracy: 0.8703125  
Validation Accuracy: 0.867025  
Train MAE: 0.1616375  
Validation MAE: 0.17395  
Train MSE: 0.489175  
Validation MSE: 0.7656

max\_depth: 12, n\_estimators:100  
Train Accuracy: 0.8752625  
Validation Accuracy: 0.8676  
Train MAE: 0.15469375  
Validation MAE: 0.173075  
Train MSE: 0.48120625  
Validation MSE: 0.726925

max\_depth: 12, n\_estimators:200  
Train Accuracy: 0.875175  
Validation Accuracy: 0.867475  
Train MAE: 0.15435625  
Validation MAE: 0.172575  
Train MSE: 0.43398125  
Validation MSE: 0.701275

max\_depth: 12, n\_estimators:400  
Train Accuracy: 0.87535  
Validation Accuracy: 0.8675  
Train MAE: 0.15424375  
Validation MAE: 0.1726  
Train MSE: 0.46011875  
Validation MSE: 0.7209

max\_depth: 12, n\_estimators:600

```
Train Accuracy: 0.87516875
Validation Accuracy: 0.867825
Train MAE: 0.15446875
Validation MAE: 0.172625
Train MSE: 0.46076875
Validation MSE: 0.702175
```

```
CPU times: user 30min 28s, sys: 10.9 s, total: 30min 39s
Wall time: 9min
```

```
In [35]: # select best hyper parameters for test data evaluation
best_params = sorted(hyper_parameter_list, key = lambda x: x[2])[0]

max_depth = best_params[0]
n_estimators = best_params[1]
```

## Evaluate Classification on Validation Data - Precision and Recall

```
In [36]: # get predictions on validation data
y_pred_val = [i for i in hyper_parameter_list if i[0] == max_depth and i[1] == n_estimators][0][3]

len(y_pred_val)
```

```
Out[36]: 40000
```

```
In [37]: # compute confusion matrix
cnf_matrix = confusion_matrix(y_val, y_pred_val)

# compute evaluation metrics
FP = cnf_matrix.sum(axis=0) - np.diag(cnf_matrix)
FN = cnf_matrix.sum(axis=1) - np.diag(cnf_matrix)
TP = np.diag(cnf_matrix)
TN = cnf_matrix.sum() - (FP + FN + TP)

# TPR, sensitivity, or recall
TPR = TP / (TP + FN)

# TNR or specificity
TNR = TN / (TN + FP)

# precision
precision = TP / (TP + FP)

# FPR
FPR = FP / (FP + TN)

# FNR
FNR = FN / (TP + FN)

# accuracy
accuracy = (TP+TN)/(TP+FP+FN+TN)
```

/Users/gio/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:11: RuntimeWarning: invalid value encountered in true\_divide  
# This is added back by InteractiveShellApp.init\_path()  
/Users/gio/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:17: RuntimeWarning: invalid value encountered in true\_divide  
/Users/gio/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:23: RuntimeWarning: invalid value encountered in true\_divide

**Table 3. Confusion matrix and classification metrics for each number of vote (outOf feature, label).**



```
In [38]: # create dataframe
columns = ['FP', 'FN', 'TP', 'TN', 'TPR', 'TNR', 'precision', 'FPR', 'FNR', 'accuracy']
rows = set(y_val).union(set(y_pred_val))
data = list(zip(FP, FN, TP, TN, TPR, TNR, precision, FPR, FNR, accuracy))
df = pd.DataFrame(data=data, index=rows, columns=columns)
df.index.names = ['Label (Number of Votes)']

# print
pd.set_option('display.max_rows', df.shape[0]+1)
df
```

Out[38]:

	FP	FN	TP	TN	TPR	TNR	precision	FPR	FNR	accuracy
Label (Number of Votes)										
0	36	1530	27595	10839	0.947468	0.996690	0.998697	0.003310	0.052532	0.960850
1	1928	625	4689	32758	0.882386	0.944416	0.708629	0.055584	0.117614	0.936175
2	863	795	1034	37308	0.565336	0.977391	0.545071	0.022609	0.434664	0.958550
3	620	387	651	38342	0.627168	0.984087	0.512195	0.015913	0.372832	0.974825
4	397	360	219	39024	0.378238	0.989929	0.355519	0.010071	0.621762	0.981075
5	182	315	89	39414	0.220297	0.995404	0.328413	0.004596	0.779703	0.987575
6	209	169	101	39521	0.374074	0.994739	0.325806	0.005261	0.625926	0.990550
7	138	161	56	39645	0.258065	0.996531	0.288660	0.003469	0.741935	0.992525
8	111	109	46	39734	0.296774	0.997214	0.292994	0.002786	0.703226	0.994500
9	82	97	29	39792	0.230159	0.997944	0.261261	0.002056	0.769841	0.995525
10	67	88	28	39817	0.241379	0.998320	0.294737	0.001680	0.758621	0.996125
11	71	59	32	39838	0.351648	0.998221	0.310680	0.001779	0.648352	0.996750
12	44	65	13	39878	0.166667	0.998898	0.228070	0.001102	0.833333	0.997275
13	49	49	14	39888	0.222222	0.998773	0.222222	0.001227	0.777778	0.997550
14	38	41	9	39912	0.180000	0.999049	0.191489	0.000951	0.820000	0.998025
15	46	34	13	39907	0.276596	0.998849	0.220339	0.001151	0.723404	0.998000
16	35	48	7	39910	0.127273	0.999124	0.166667	0.000876	0.872727	0.997925
17	28	22	10	39940	0.312500	0.999299	0.263158	0.000701	0.687500	0.998750

	FP	FN	TP	TN	TPR	TNR	precision	FPR	FNR	accuracy
Label (Number of Votes)										
18	25	27	6	39942	0.181818	0.999374	0.193548	0.000626	0.818182	0.998700
19	28	16	8	39948	0.333333	0.999300	0.222222	0.000700	0.666667	0.998900
20	15	17	4	39964	0.190476	0.999625	0.210526	0.000375	0.809524	0.999200
21	18	18	5	39959	0.217391	0.999550	0.217391	0.000450	0.782609	0.999100
22	22	13	6	39959	0.315789	0.999450	0.214286	0.000550	0.684211	0.999125
23	15	15	3	39967	0.166667	0.999625	0.166667	0.000375	0.833333	0.999250
24	13	14	1	39972	0.066667	0.999675	0.071429	0.000325	0.933333	0.999325
25	15	17	4	39964	0.190476	0.999625	0.210526	0.000375	0.809524	0.999200
26	6	12	1	39981	0.076923	0.999850	0.142857	0.000150	0.923077	0.999550
27	14	10	3	39973	0.230769	0.999650	0.176471	0.000350	0.769231	0.999400
28	11	10	3	39976	0.230769	0.999725	0.214286	0.000275	0.769231	0.999475
29	14	9	2	39975	0.181818	0.999650	0.125000	0.000350	0.818182	0.999425
30	9	12	0	39979	0.000000	0.999775	0.000000	0.000225	1.000000	0.999475
31	7	9	0	39984	0.000000	0.999825	0.000000	0.000175	1.000000	0.999600
32	8	9	3	39980	0.250000	0.999800	0.272727	0.000200	0.750000	0.999575
33	4	11	0	39985	0.000000	0.999900	0.000000	0.000100	1.000000	0.999625
34	4	5	1	39990	0.166667	0.999900	0.200000	0.000100	0.833333	0.999775
35	5	2	2	39991	0.500000	0.999875	0.285714	0.000125	0.500000	0.999825
36	12	2	3	39983	0.600000	0.999700	0.200000	0.000300	0.400000	0.999650
37	4	9	0	39987	0.000000	0.999900	0.000000	0.000100	1.000000	0.999675
38	7	8	1	39984	0.111111	0.999825	0.125000	0.000175	0.888889	0.999625
39	2	3	0	39995	0.000000	0.999950	0.000000	0.000050	1.000000	0.999875
40	3	3	0	39994	0.000000	0.999925	0.000000	0.000075	1.000000	0.999850
41	2	4	1	39993	0.200000	0.999950	0.333333	0.000050	0.800000	0.999850
42	4	2	0	39994	0.000000	0.999900	0.000000	0.000100	1.000000	0.999850
43	4	1	0	39995	0.000000	0.999900	0.000000	0.000100	1.000000	0.999875

	FP	FN	TP	TN	TPR	TNR	precision	FPR	FNR	accuracy
Label (Number of Votes)										
44	4	5	0	39991	0.000000	0.999900	0.000000	0.000100	1.000000	0.999775
45	8	3	1	39988	0.250000	0.999800	0.111111	0.000200	0.750000	0.999725
46	4	4	1	39991	0.200000	0.999900	0.200000	0.000100	0.800000	0.999800
47	2	4	0	39994	0.000000	0.999950	0.000000	0.000050	1.000000	0.999850
48	3	4	1	39992	0.200000	0.999925	0.250000	0.000075	0.800000	0.999825
49	1	3	0	39996	0.000000	0.999975	0.000000	0.000025	1.000000	0.999900
50	0	3	1	39996	0.250000	1.000000	1.000000	0.000000	0.750000	0.999925
51	1	2	0	39997	0.000000	0.999975	0.000000	0.000025	1.000000	0.999925
52	7	3	0	39990	0.000000	0.999825	0.000000	0.000175	1.000000	0.999750
53	4	3	0	39993	0.000000	0.999900	0.000000	0.000100	1.000000	0.999825
54	2	3	0	39995	0.000000	0.999950	0.000000	0.000050	1.000000	0.999875
55	3	4	0	39993	0.000000	0.999925	0.000000	0.000075	1.000000	0.999825
56	2	2	0	39996	0.000000	0.999950	0.000000	0.000050	1.000000	0.999900
57	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
58	2	2	1	39995	0.333333	0.999950	0.333333	0.000050	0.666667	0.999900
59	2	1	0	39997	0.000000	0.999950	0.000000	0.000050	1.000000	0.999925
60	2	0	0	39998	NaN	0.999950	0.000000	0.000050	NaN	0.999950
62	1	3	0	39996	0.000000	0.999975	0.000000	0.000025	1.000000	0.999900
63	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
64	0	4	0	39996	0.000000	1.000000	NaN	0.000000	1.000000	0.999900
65	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
66	1	2	0	39997	0.000000	0.999975	0.000000	0.000025	1.000000	0.999925
67	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
68	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
69	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
70	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975

	FP	FN	TP	TN	TPR	TNR	precision	FPR	FNR	accuracy
Label (Number of Votes)										
71	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
72	1	2	0	39997	0.000000	0.999975	0.000000	0.000025	1.000000	0.999925
73	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
74	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
75	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
76	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
77	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
79	2	0	0	39998	NaN	0.999950	0.000000	0.000050	NaN	0.999950
81	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
83	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
84	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
86	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
87	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
88	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
90	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
92	0	3	0	39997	0.000000	1.000000	NaN	0.000000	1.000000	0.999925
93	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
94	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
96	0	0	1	39999	1.000000	1.000000	1.000000	0.000000	0.000000	1.000000
97	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
98	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
102	0	0	1	39999	1.000000	1.000000	1.000000	0.000000	0.000000	1.000000
104	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
105	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
106	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
108	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950

	FP	FN	TP	TN	TPR	TNR	precision	FPR	FNR	accuracy
Label (Number of Votes)										
109	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
112	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
113	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
114	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
116	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
117	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
118	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
125	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
130	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
135	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
138	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
160	1	1	0	39998	0.000000	0.999975	0.000000	0.000025	1.000000	0.999950
161	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
165	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
182	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
187	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
204	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
209	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
221	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
227	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
231	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
236	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
242	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
286	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975
377	1	0	0	39999	NaN	0.999975	0.000000	0.000025	NaN	0.999975
384	0	1	0	39999	0.000000	1.000000	NaN	0.000000	1.000000	0.999975

**Figure 2. Precision and recall for random forest model implementation (max\_depth: 12, n\_estimators: 200, number of features: 22).**

In [39]: *# plot precision and recall*

```
plt.figure(figsize=(15,7)) # define figure
```

```
# precision
```

```
plt.scatter(list(rows), precision, facecolors='none', edgecolors='b')
```

```
# recall
```

```
plt.scatter(list(rows), TPR, facecolors='none', edgecolors='r')
```

```
# add labels
```

```
plt.legend(['Precision', 'Recall'], fontsize=14)
```

```
plt.title('Precision and Recall for Random Forest Model on Validation Dataset', fontsize=14)
```

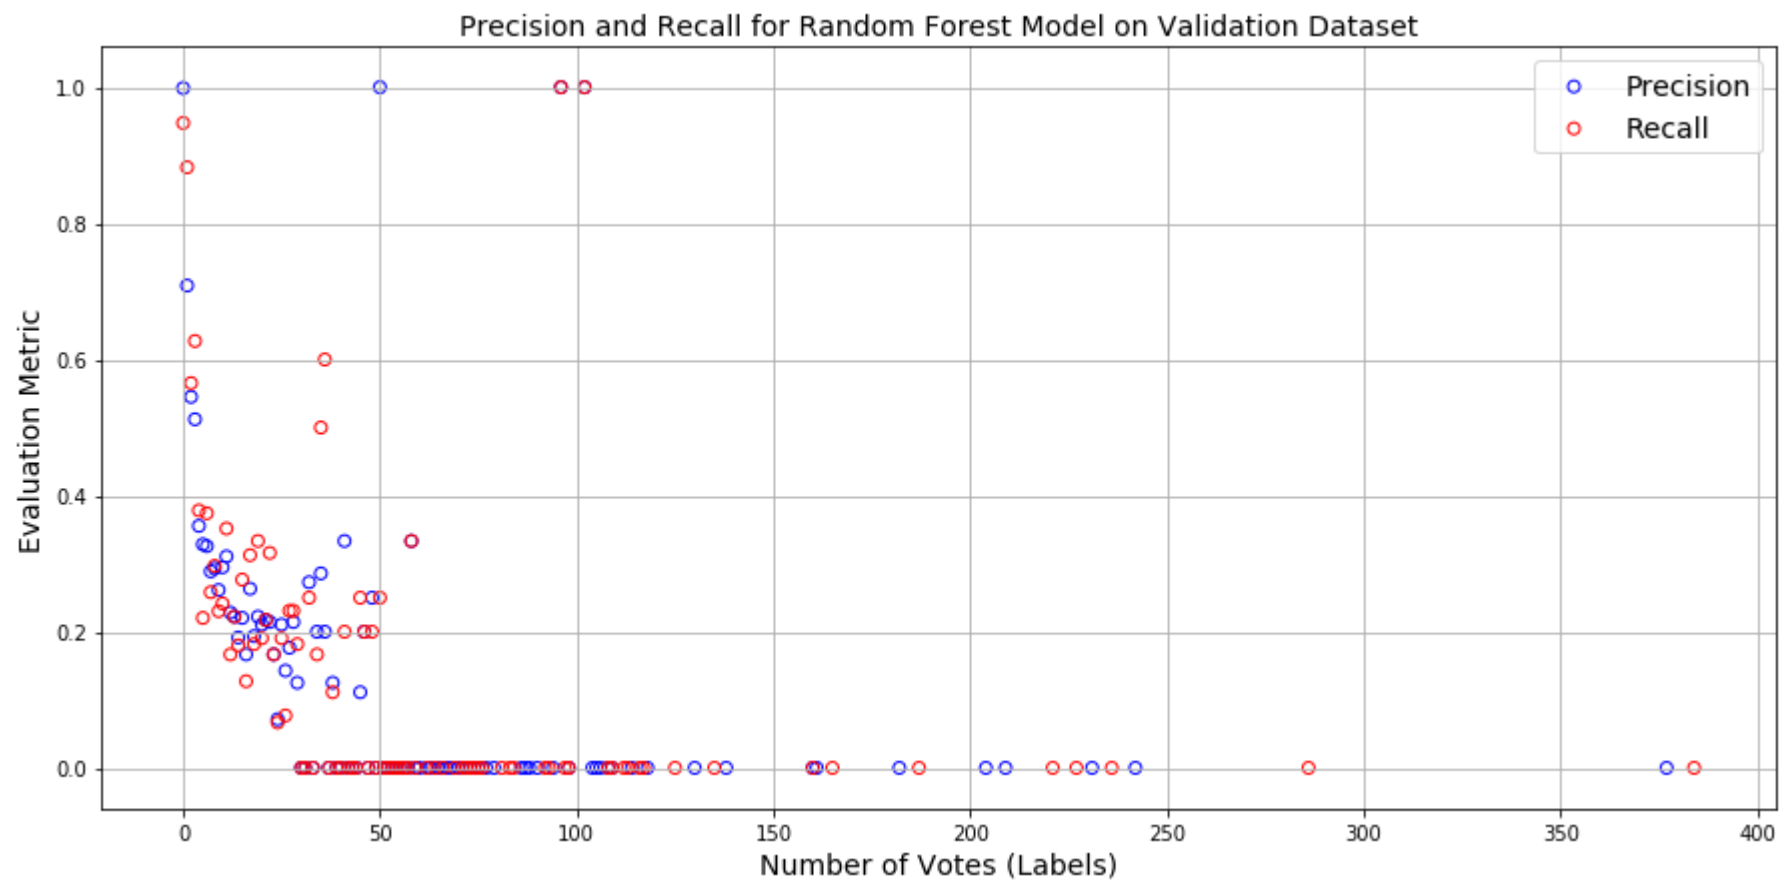
```
plt.xlabel('Number of Votes (Labels)', fontsize=14)
```

```
plt.ylabel('Evaluation Metric', fontsize=14)
```

```
# show plot
```

```
plt.grid()
```

```
plt.show()
```



Evaluate on Test



In [40]: %%time

```
# get indices
idx = np.array([i['outOf'] for i in train_features.helpful]) > 0

# create features
X_train_val = Get_features(train_features)
X_test = Get_features(test)

# create labels
y_train_val = [i['nHelpful'] for i in train_labels]
y_test = None

# find filtered test data
X_train_val_filt = Get_features(train_features.loc[idx,:])
y_train_val_ratio_filt = Get_labels_ratio(list(compress(train_labels, idx)))
```

CPU times: user 4min 4s, sys: 4.61 s, total: 4min 8s  
Wall time: 4min 11s

In [41]: %%time

```
# (train_accuracy, train_mae, train_mse, y_pred_test)
metrics_tuple = perfmetrics_RFmodel(max_depth, n_estimators, X_train_val_filt[features_tokeep],
                                     y_train_val_ratio_filt, X_train_val[features_tokeep],
                                     y_train_val, X_test[features_tokeep], y_test)

# print metrics
print('max_depth: {}, n_estimators:{}'.format(max_depth, n_estimators))
print('Train Accuracy: {}'.format(metrics_tuple[0]))
print('Train MAE: {}'.format(metrics_tuple[1]))
print('Train MSE: {}'.format(metrics_tuple[2]))
print('\n')
```

max\_depth: 12, n\_estimators:200  
Train Accuracy: 0.875125  
Train MAE: 0.15479  
Train MSE: 0.41986

CPU times: user 3min 18s, sys: 1.59 s, total: 3min 20s  
Wall time: 1min 5s

```
In [42]: # save predictions on file "predictions_Helpful.txt"
save_predictions(metrics_tuple[3])
```

```
In [ ]:
```